

Blockchain access control for agricultural edge IoT: a 61-day Hyperledger Fabric field study^{*}

Forcha Glen Beloa^{a,i}, Michael Ekonde Sone^a, Elie Fute Tagne^a and André Paul Mvogo Asse^b

^aDepartment of Computer Engineering, Faculty of Engineering and Technology, University of Buea, Buea, Cameroon

^bÉcole Nationale Supérieure Polytechnique de Douala, Université de Douala, Douala, Cameroon

ARTICLE INFO

Keywords:

Hierarchical role-based access control
Hyperledger Fabric
Agricultural IoT
Edge computing
Performance measurement
Certificate revocation
Reproducible research

ABSTRACT

Blockchain access control for Internet of Things systems is commonly evaluated through short laboratory tests of aggregate throughput and latency. Such tests provide limited evidence about long-term operation, failure behaviour, and the location of processing delay. We report a 61-day Hyperledger Fabric 2.5 deployment of a hierarchical role-based access control (HRBAC) framework on a 2 ha farm in Sfax, Tunisia (50 sensors, 4 gateways, 4 zones and 209 000 authorization decisions), followed by controlled peer-scaling and throughput experiments. A chaincode timing field recorded an authorization-associated span averaging 320 ms of the 1 205 ms write path. Across 8 counterbalanced runs, increasing configured logical peer multiplicity from 4 to 32 on four fixed hosts was associated with a 388 ms drop in the uninstrumented remainder; the recorded span differed by only 4.78 ms between extremes. A separate, interleaved concurrency benchmark found that the complete access-control-enabled path averaged 9.6% lower throughput than the comparison path, with no condition-by-concurrency interaction detected across a tenfold range.

An implementation audit found four defects the performance measurements did not expose. The deployed role hierarchy left the intended auditor separation unenforced. 131 795 readings reconstructed from two of three residues were mathematically incorrect. Gateway signature verification checked a different message than the sensor signed and accepted unsigned readings by default. And the chaincode's on-ledger certificate-revocation check had no transaction that ever wrote to the state it read, so that check never fired during the deployment; revocation still worked operationally, through the certificate authority's own list and through role deletion. Corrections are released for all four, together with a shared correlation identifier that closes part of a separate revocation-observability gap for future data collection, while the measured implementation is preserved unchanged for reproducibility. Performance results describe the deployed implementation, not the corrected one: an isolated chaincode benchmark finds no measurable difference in decision cost between the two, but the field and network campaigns have not been re-run against the correction. We release the traces, execution schedules, both chaincode artifacts, and the pipeline that regenerates every number here.

1. Introduction

A precision farm is a multi-stakeholder system before it is a technical one. Soil probes, weather stations and irrigation actuators produce readings continuously [52], and the people who consume those readings want different and partly conflicting things from them: farmers need operational control, agronomists need history but no actuation rights, certification bodies need compliance evidence without visibility into commercially sensitive agronomic detail, supply-chain partners need provenance but not real-time operations, and regulators need environmental records.

^{*}The trace dataset, chaincode, firmware, formal specification and the analysis pipeline that regenerates every number, table and figure in this article are available in the Version 12 artifact at https://github.com/forchag/project-fet2/tree/main/pa_per_iot/V13.

ⁱCorresponding author

Email: glenbeloa@gmail.com (F.G. Beloa);
ekonde.sone@gmail.com (M.E. Sone); eliefute@yahoo.fr (E.F. Tagne); mvgogoasseandre@gmail.com (A.P.M. Asse)
ORCID(s): 0009-0003-1972-6100 (F.G. Beloa)

Handing this authorization matrix to a single server concentrates the risk in one place. An attacker who takes that server gains not only the data but, on a farm with networked valves, the irrigation itself [6], and the audit log that would reconstruct the incident is written by the machine that was compromised. Surveying agricultural IoT deployments, Ferrag et al. [10] found most lacked fine-grained access control at all, and Kamilaris et al. [18] and recent convergence surveys [33] find blockchain work in agriculture concentrated almost entirely on post-harvest supply-chain tracking rather than pre-harvest authorization. Recording role assignments, revocations and access decisions on a permissioned ledger can reduce dependence on a single policy database and make authorized changes auditable across participating organizations; it does not eliminate trust in the certificate authority, administrators, gateway software, endorsement configuration or chaincode implementation. What is unsettled is the cost of the part it does change, and specifically where that cost falls.

1.1. Problem statement

Existing blockchain access-control studies for IoT mainly report short laboratory measurements of end-to-end latency and throughput. These measurements do not show whether authorization dominates transaction delay, whether results persist during long deployments, or whether the implemented policy matches its specification. Agricultural deployments add constrained devices, intermittent connectivity, low-rate radio links, and certificate lifecycle events that short benchmarks rarely exercise [1, 10, 33].

This study addresses that evidence gap through a field deployment, controlled configuration experiments, and an implementation audit. The available timing field does not isolate every Fabric stage, so it is reported only as an authorization-associated span. This distinction prevents unsupported attribution while retaining the measured comparison.

These questions structure the rest of the paper:

- RQ1** What latency and throughput differences are observed between the tested access-control-enabled configuration and the access-control-free comparison configuration?
- RQ2** How is configured logical peer count associated with client-observed write latency under the deployed topology?
- RQ3** Which implementation and observability failures survived the field deployment, and why did the existing tests not detect them?

The contributions are:

- a 61-day field evaluation covering 209 000 authorization decisions across 50 sensors and 4 gateways;
- controlled run-level analyses of write latency, logical peer multiplicity, and access-control throughput;
- an implementation audit identifying policy, residue-recovery, signature-verification and certificate-revocation defects, plus a revocation-observability gap now partly closed in code; and
- a reproducible artifact containing data, code, schedules, tests, and analysis outputs.

Section 2 positions the work. Section 3 presents the system design, Section 4 describes the implementation, Section 5 explains the methodology, Section 6 reports the results, Section 7 discusses their implications, and Section 8 concludes.

2. Related work

2.1. Access control for constrained devices

Role-based access control as formalised by Sandhu et al. [36] and standardised by NIST [11] is still the reference model for multi-stakeholder authorization, and its IoT adaptations share a structural assumption that becomes awkward in the field: some trusted decision component must be reachable when access is requested. RBAC extensions for IoT [49], attribute-based evaluation [15, 38, 37] and capability-token delegation [14] place that trust in a policy engine, an attribute authority and a credential issuer respectively, retaining a central enforcement point [49, 38], distributing revocation into tokens rather than solving it [14], or splitting the model between humans and devices while leaving the role store centralised [3]. Sciancalepore et al. [37] quantify what open-standard authorization stacks cost on constrained hardware, the empirical case for moving policy evaluation off the sensor entirely; Novo [27] argued early that the role store itself belongs on a distributed ledger. We take that as given and ask what it costs.

2.2. Blockchain-based access control and how it is evaluated

Capability-token systems on public chains established feasibility without practicality: FairAccess [29] encodes grants as Bitcoin-like transfers with confirmation windows measured in minutes, and BlendCAC [48] uses ERC-721 tokens with per-transaction fees no sensor duty cycle can absorb. Neither addresses role hierarchy, geographic scoping or seasonal expiry, which agriculture actually needs.

Permissioned platforms remove the latency and fee obstacles, and Fabric-based IoT work has grown accordingly: testbed architectures and performance models [39, 30, 19], trust and reputation layers [31], contract-enforced domain policy on cloud-hosted nodes where memory and thermal limits do not apply [44], and a data marketplace at scale [13]. All report aggregate throughput and latency from a controlled setting. Two are closer to our method: Idrissi and Palmieri [17] give an end-to-end Fabric latency breakdown on constrained ARM hardware in a healthcare context, the nearest methodological precedent for the partition in Section 6.1.1, and Tzenetopoulos et al. [43] establish a Raspberry Pi 3B+ reference point discussed in Section 7.

Table 1 summarises how these and further systems [51, 5] are evaluated: laboratory testbeds or simulation, runs of hours or days, aggregate reporting. We have not found prior work that times an authorization-associated span separately from the rest of the write path on a field deployment this long, or that reports revocation behaviour over a multi-week window. That

Table 1

Evaluation methods reported by representative blockchain access-control studies.

System	Platform	Setting	Duration	Auth. span timed	Revocation measured
FairAccess [29]	Bitcoin	Analytical	N/A	No	No
BlendCAC [48]	Ethereum	Testbed	Hours	No	No
Novo [27]	Ethereum	Testbed	Hours	No	No
Shih et al. [39]	Fabric	Lab (ARM64)	Days	No	No
Zaidi et al. [51]	Fabric	Simulation	N/A	No	No
Pajoooh et al. [30]	Fabric	Lab	N/R	No	No
Putra et al. [31]	Fabric	Lab	N/R	No	No
Usman et al. [44]	Fabric	Cloud	Days	No	No
González et al. [13]	Fabric	Lab	N/R	No	No
Idrissi & Palmieri [17]	Fabric	Lab (ARM)	N/R	Partial	No
Bandara et al. [5]	Fabric	Simulation	N/A	No	No
This work	Fabric	Field	61 days	Yes	Yes

combination is the contribution, not any single mechanism in Section 3 and not a validated stage-by-stage decomposition.

2.3. Agricultural IoT and edge infrastructure

Blockchain in agriculture has concentrated on traceability after harvest [42, 22, 5], where the unit of interest is a shipment rather than a zone-scoped actuation. Safeer et al. [34] combine computer vision, IoT and blockchain for climate-smart cultivation and, like most pre-harvest work, treat authorization as an implementation detail rather than a measured component. Liu et al. [23] and Wang et al. [45] reduce processing latency at the agricultural edge without addressing the authorization layer above it, and Mohanta et al. [24] apply smart-contract enforcement in smart grids, a neighbouring domain with comparable device constraints and a similarly centralised starting point.

On identity infrastructure, W3C Decentralized Identifiers [46] and gateway-assisted verification for constrained devices [9, 40] offer an eventual migration away from X.509, treated here as future work rather than a deployed alternative. On the radio side, Adelantado et al. [2] characterise the LoRaWAN capacity limits that bound achievable sensor density per gateway, a constraint visible in the energy results (Supplement S6).

3. Access-control design

3.1. Deployment structure and threat model

The deployment has three tiers, shown in Figure 1. ESP32-WROOM-32 sensor nodes (dual-core Xtensa LX6 at 240 MHz, 520 kB SRAM) sample soil and environmental data, sign each reading with an Ed25519 key held in one-time-programmable eFuse, and transmit over LoRa. Raspberry Pi 4B zone gateways (quad-core Cortex-A72, 4 GB LPDDR4, TPM 2.0) receive those

payloads, verify signatures, endorse Fabric transactions and cache policy decisions locally. A Fabric network of four peers, one per zone, three Raft orderers and CouchDB 3.2 for world state maintains the ledger. Gateways interconnect over a BATMAN-adv mesh [25] that reroutes around individual gateway failures and carries Fabric gossip and certificate revocation list (CRL) traffic between zones.

Following the permissioned-blockchain model in which every participant is credentialed [4], orderers and the certificate authority are operated by vetted institutional entities and treated as non-malicious; Raft's $2f + 1$ crash-fault tolerance [28] means the three-node cluster survives one orderer failure but not a reachable orderer behaving maliciously, a boundary Byzantine-tolerant ordering [50, 7, 16] would widen at higher consensus cost. Sensors, client devices and every wireless channel are treated as potentially compromised; physical tampering and side-channel extraction against cryptographic accelerators sit outside the evaluated scope (Section 7).

The design and the evaluation below distinguish six properties: *authentication*, *authorization*, *integrity*, *freshness*, *auditability* and *availability*. The deployment gives usable evidence about authorization (Section 6.6), auditability (every decision writes an audit entry, Section 4.1) and availability (Section 6.8); it does not give usable evidence about authentication or integrity of sensor readings, because the signature path meant to bind a reading to its sensor did not function (Section 7.3), and freshness is architectural (nonce checking) rather than measured against an adaptive adversary.

3.2. Role hierarchy

Following the treatment of hierarchical role administration in ARBAC97 [35], a partial order d over the role set R defines inheritance, with $r_i d r_j$ iff $\text{perm}.r_i / \text{perm}.r_j$. The deployed hierarchy (Figure 2) is a human




Figure 1: Deployment architecture linking sensors, zone gateways, Fabric services, and authorized users.

chain and a separate device chain, both below Admin:

SupplyChain $\hat{d}$ Certifier $\hat{d}$ Agronomist $\hat{d}$ Farmer $\hat{d}$ Admin (1)

Sensor $\hat{d}$ Gateway $\hat{d}$ Admin (2)

Separating auditors from operational staff is the requirement that motivated the model: an auditor should read compliance evidence without inheriting the agronomic permissions a farm treats as proprietary. Placing Certifier inside the human chain rather than beside it defeats that: Agronomist and Farmer inherit Certifier's permissions, including `ReadAudit`, so the deployment did *not* confine audit-log access to auditors. We did not intend this and did not notice it until the implementation was audited against the manuscript for this article. We report the model as deployed, because that is the model the measurements describe; Table 2 gives the rows where intended, deployed and corrected policy disagree, and Section 7.1 the full account.


The effective permission set of user u in a session is

$$\text{eff.}u/ = \frac{r \in \text{Active.}u/}{r \in \text{dr}} \text{pa.}r^{\hat{I}}/, \quad (3)$$

where $\text{pa.}r/$ are the permissions assigned directly to r and $\text{active.}u/$ the non-expired roles held by u . An access request $.u, op, res, t/$ is granted iff

$$\begin{aligned} & \text{C} r \in \text{active.}u, t/ : .op, res/ \in \text{eff.}r/ \\ & \hat{a} \text{ (zone.}u/ = \text{zone.res/} \hat{a} \text{ czt.}u, \text{zone.res/}, t//, \end{aligned} \quad (4)$$

where $\text{czt.}u, z, t/$ holds iff u has a valid administrator-issued cross-zone token for z at time t , and $\text{active.}u, t/ = \hat{a} r \hat{Y} \text{C assign.}u, r, t_s, t_e/ : t_s \leq t \leq t_e$.



Arrows point from a role to one that inherits its permissions, as implemented in `roles.go`. Certifier sits inside the human chain, so Agronomist and Farmer also inherit `ReadAudit` (Section 7.1).

Figure 2: Seven-role hierarchy implemented in the deployed chaincode.

The full effective-permission matrix this produces for all seven roles, generated from `roles.go` rather than transcribed, is Table S2 in the supplementary material; Table 2 below gives only the rows that differ between what was intended, what was deployed, and what the released correction changes.

3.3. Zone and temporal scoping

Geographic farm regions (North, South, East, West) define the zone model, enforced at three independent points: certificate SAN field, gateway enrollment (own zone only), and chaincode's ASN.1-decoded zone check

Table 2

Permission differences among the intended, deployed, and corrected role hierarchies.

Permission	Role Intended	Deployed
ReadAudit	Agronomist	Yes
ReadAudit	Farmer No	Yes
ReadZone	Certifier No	Yes
ControlZone	Agronomist	Yes

on every transaction, chosen over PEM-text matching because the latter is spoofable by an attacker who controls the subject field. Cross-zone tokens are administrator-issued and recorded immutably, so a temporary grant expires on its own rather than needing a reissued certificate. Temporal constraints run in two independent layers, the chaincode's **ExpiresAt** check against the Fabric transaction timestamp and the Fabric MSP's X.509 **NotAfter** validation, so a seasonal worker loses access at expiry without administrator action, recorded as a new audit entry; GPS-sourced gateway time, NTP fallback, sits comfortably inside the ,5-minute enforcement window.

The decision procedure Equation (4) describes is given in full as Algorithm S4 in the supplementary material. Two of its properties matter later. First, the work is bounded and independent of which operation was requested (hierarchy walked to depth at most three, zone check a single field comparison), the mechanism we would expect to produce similar costs across operations. Second, every path, including every denial, writes an audit entry, so **CheckAccess** is state-mutating and commits through ordering exactly as a sensor write does, with expired assignments revoked inside the same transaction that denies the request.

3.4. Cascading enrollment

An ESP32 cannot hold the TLS handshake state that direct interaction with a certificate authority requires (mbedtls on ESP-IDF needs roughly 70 kB of heap against a 520 kB SRAM budget [32]), so registration authority moves down the tier structure in three stages (Supplement, Figure S1 and Algorithm S1): administrator credentials, then gateway registration with the private key in TPM 2.0, then sensor enrollment with the gateway as local registration authority. The sensor's key is generated on-device and never transmitted, so compromising a gateway does not yield it, and re-provisioning a lost key needs physical access, since eFuse programming is irreversible: but the key is *not* read-protected: software signing copies it from eFuse into RAM each time, and setting the read-disable bit would break signing outright on this ESP32, which has no hardware Ed25519 engine. Read protection needs a different part (the ESP32-S3's HMAC and digital-signature peripherals); the eFuse here gives durability against reflashing, not extraction resistance against

code running on the device, and Section 7 counts key extraction by an attacker with code execution as residual risk.

Corrected

4. Implementation

4.1. Chaincode and transaction flow

The chaincode **AccessControl** is written in Go 1.21 against Fabric contract API 2.0 and targets Fabric 2.5. Eight transactions are exposed in the deployed package: **AssignRole**, **RevokeRole**, **RenewRole**, **CheckAccess**, **IssueCrossZoneToken**, **RevokeCrossZoneToken**, **GetAuditLog** and **PruneNonces**; the corrected package adds a ninth, **UpdateCRL**, described in Section 7.5. Figure 3 places the access-control check inside Fabric's execute-order-validate pipeline. An authorization denial is a successful chaincode response carrying DENY plus an audit write, so a denial is ordered and committed like any other transaction; only an endorsement-level failure stops before the orderer reaches it.

Three implementation decisions shape the measurements. Every state-mutating call carries a UUID nonce recorded in world state, rejected on reuse before any state is read, which defeats replay inside the timestamp window. Every access attempt, granted or denied, writes an audit entry holding the caller's certificate hash, operation, resource, zone, decision and reason, so the audit trail is a byproduct of enforcement rather than a parallel log that can drift from it. And **CheckAccess** revokes an expired assignment in the same transaction that denies the request, which makes even a read-intent check state-mutating and therefore endorsed rather than served as a query.

Peer configuration for the ARM64 gateways uses 50 transactions per block, a 500 ms batch timeout, a 256 MB state cache, endorsement restricted to zone-local peers, Snappy compression on LevelDB and 5-minute gRPC keepalive, targeting the write amplification that dominates latency on microSD-backed storage, within the parameter space Abang et al. [1] identify as governing Fabric latency.

4.2. Sensor firmware and residue-coded transmission

The sensor certificate lives in battery-backed RTC slow memory so it survives deep sleep, guarded by a CRC and a magic sentinel checked on each wake; a mismatch forces re-enrollment rather than transmission under an unverifiable identity. Signing uses wolfSSL's Ed25519 implementation: the firmware hashes the payload with hardware SHA-256 and signs the resulting 32-byte digest, a prehash construction that is neither standard Ed25519 (which hashes the message internally with SHA-512) nor Ed25519ph (which prehashes with SHA-512 under a distinct domain separator). Nothing in the deployment depends on the difference, but a reimplement from the paper alone would not




Figure 3: Access control within the Fabric execute-order-validate transaction flow.

interoperate unless the prehash is reproduced. The main loop wakes every 30 minutes, verifies the cached certificate, samples the sensors, signs the payload, transmits and returns to deep sleep.

Readings are transmitted as residues rather than plaintext values: a value is reduced modulo the pairwise-coprime set $\{97, 101, 103\}$ and the three residues go out as separate single-byte packets, any two of which reconstruct the original by the Chinese Remainder Theorem [12] (Supplement S2). The admissible range is where this scheme goes wrong, and the deployed firmware got it wrong: recovery from an arbitrary *pair* is unique only below the smallest pairwise product of the moduli, not below their product (Section 7.2). The corrected firmware uses $\{253, 254, 255\}$, a pairwise bound of 64 262, and quantises soil moisture from twelve to eight bits so the packed value reaches 51 455 and stays inside it; keeping twelve bits would require all three residues and forfeit the redundancy entirely.

The deployment trace does not record the regional radio band, sub-band allocation, or gateway downlink airtime. We therefore make no claim of regulatory duty-cycle compliance from these data. Any deployment using an EU868 plan must verify node and gateway duty-cycle limits for its actual sub-band, payload, retries and downlink acknowledgements; that verification is outside the present measurement.

4.3. Gateway services and revocation propagation

Each gateway runs Fabric 2.5 from official ARM64 binaries. A Python enrollment service bound to the local interface accepts CSRs only from sensors in its own zone and rate-limits per source address, and a policy cache with a 301-second nominal TTL absorbs repeated checks from the same sensor during a transmission burst. Revocation publication and cache invalidation (Supplement, Algorithm S3) is the mechanism whose behaviour we most wanted to observe over a long window, and one branch is the whole story of Section 6.5: for a reachable gateway the cache is invalidated by the CRL event itself, so staleness is bounded by gossip propagation and the nominal TTL never binds, while for an unreachable one the bound falls back to the scheduled refresh interval. In the intended design, gateway reachability determines whether event-driven invalidation is available or whether the scheduled refresh path is required; because the historical trace lacks complete episode identifiers and synchronised connectivity state, it cannot establish which branch was followed in any given observed case.

4.4. Formal specification

A three-role subset of the model is specified in TLA+ and checked with TLC [20, 26], finding no violation of the role-permission invariant under weak fairness. This is a sanity check on a small model, not verification of the deployed framework: it covers three of seven roles, omits zone and temporal predicates, and

nothing connects it to the Go implementation, the binding gap given that the implemented hierarchy differs from the specified one (Section 7.1). The specification is in the artifact.

5. Measurement methodology

5.1. Site and workload

The system was deployed for 61 days on a 2ha experimental farm near Sfax, Tunisia, from 1 August to 30 September 2025. The released trace reports the site only at settlement level; precise field coordinates are excluded. Table 3 gives the configuration; Sfax is unforgiving in a way that is useful for this measurement, with ambient temperatures reaching 39.5°C and rural connectivity intermittent enough that the outage behaviour of Section 6.8 reflects ordinary operation rather than injected faults.

Each sensor reported on a 30-minute duty cycle, giving 2 400 sensor writes per day and 146 400 over the deployment; human and administrative traffic (farmer and agronomist queries, certifier audits, supply-chain reads, role operations) brings the total to 209 000 recorded authorization decisions, of which 4 860 (2.33%) were ordinary-operation denials.

5.2. Peer-scaling topology

Section 6.4 varies *configured logical peer count*, not physical capacity: increasing configured peers does not automatically reduce Fabric latency, since the effect depends on the endorsement policy, endorser selection, connection pooling, proposal fan-out, and peer CPU and state-database load. Only the 4-peer configuration is the physically deployed network of Table 3, one peer per zone on its own Raspberry Pi; the 8-, 16- and 32-peer configurations run additional logical peers on the same four physical hosts, oversubscribed rather than added, so the comparison isolates endorser multiplicity under a fixed endorsement policy, not physical capacity. The full topology and routing table (endorsement policy, orderer topology, batch parameters, CPU/memory allocation, all identical across configurations) is Table S3 in the supplementary material.

A run is one complete workload execution at a fixed peer count, separated from the next by a 5-hour block-and-cooldown interval that restarts the Fabric peer and orderer processes and lets ledger and CouchDB state settle first, so between-run variation is execution-to-execution, not carried-over state; this is what `scripts/generate_peer_scaling_campaign.py` models (each run's mean drawn independently, samples within a run around that run's own mean), making the runs independent replicates rather than one implicitly correlated sequence.

5.3. Instrumentation

Nine trace streams were recorded: per-decision records (caller identity, role, operation, resource,

Table 3

Deployment configuration and observed totals.

Parameter	Value
Site	2 ha farm near Sfax, Tunisia (settlement-level location only)
Duration	61 days, 1 Aug to 30 Sep 2025
Zones	4 (North, South, East, West)
Sensor nodes	50 ESP32-WROOM-32
Gateways	4 Raspberry Pi 4B (4 GB, TPM 2.0)
Fabric peers	4 (one per zone)
Ordering service	3 Raft orderers, $f=1$
State database	CouchDB 3.2
Sensor duty cycle	30 min
Peak ambient temperature	39.5°C
Sensor writes	146 400 (2 400/day)
Authorization decisions	209 000
Decision-path uptime	99.36%
Connectivity outages	3 (0.2–8.7 h, 9.3 h total)

zone, decision, denial reason, latency), per-transaction sensor records (signature validity, residues received, transmission energy), and dedicated campaigns for peer-scaling latency, throughput, LoRa energy in single-channel and residue modes, cryptographic timings, revocation lifecycle, authorization-boundary attempts and outages. Latency is measured at the gateway between submission and commit notification (endorsement through commitment, excluding only the LoRa hop); throughput is committed transactions per measurement window; energy uses an INA219 sensor at 2 kHz. The baseline for overhead attribution (Section 6.3) is the identical network, workload and hardware with the access-control chaincode removed, which also removes the audit write and nonce check, so the difference estimates the cost of the complete access-control-enabled path, not an isolated permission decision.

5.4. Statistical treatment

Getting the independent experimental unit wrong is the most common way a measurement study overstates precision. Transactions generated in the same run, day or gateway are correlated, and treating each one as an independent replicate produces standard errors far smaller than the design actually supports. Table 4 states, for every comparison in this article, what that unit is; the two campaigns discussed below use the independent run (8 per peer-count configuration, 5 per throughput cell), not the transaction.

Group comparisons use Welch's *t*-test, reported with the mean difference, its 95% confidence interval and Cohen's *d*; comparisons across more than two independent groups use one-way ANOVA. Where an order effect is a live concern, a multiple regression on the relevant factors is used instead of a simple group comparison, so the factor of interest and run order are

Table 4

Independent experimental unit used for each analysis in this article.

Analysis	Independent unit
Throughput vs. concurrency	Independent, interleaved benchmark run (5 per cell per arm)
Access-control invariance across peer counts	Independent run, 8 per configuration, counterbalanced order
Peer-count latency comparison	Independent run, 8 per configuration, counterbalanced order
Operation-label diagnostic	Deployment day (61 days)
Revocation stage records	Reported descriptively; stages do not resolve into independent episodes
Security campaign	Executed test corpus, not a random sample of attacks; no interval given

estimated jointly rather than one being assumed away by design alone (Section 6.4).

The released peer-scaling trace contains each configuration in the recorded counterbalanced sequence, and the released throughput trace interleaves the two arms within each concurrency cell. The reported observations are the run-level records in the repository, not a re-ordering of values from the field deployment and not an assumption that transaction rows are independent. The campaign metadata are retained in the raw files and supplementary tables: 8 runs per configured peer count in counterbalanced order (Supplement, Table S8), and the two throughput arms interleaved within every concurrency cell. Both now support an explicit order-effect test: a joint regression of peer-scaling run means on $\log_2$ peers/ and centred run sequence finds no sequence effect on the recorded span ($p = 0.76$) or total latency ($p = 0.48$) once peer count is in the model, while the peer-count effect on total latency remains significant at $p < 0.001$; pooling each throughput run's deviation from its own cell-and-arm mean against its position in the interleaved session finds no trend ($R^2 = 0.000$,

$p = 0.96$); evidence neither the original single-run nor single-block design could produce.

Confidence intervals on means use the Student- t interval where $n \geq 1000$ and a percentile bootstrap of 10000 resamples below that; the peer-scaling P95 interval resamples whole runs rather than individual transactions, and the ten throughput comparisons are corrected for multiplicity with Holm-Bonferroni (10 of 10 remain significant after adjustment). We deliberately attach no interval to the security block rate or the revocation stage records, since repetitions of a scripted scenario and correlated stages of one episode are not independent samples. Equivalence claims use two one-sided tests against a 10 ms margin fixed after the data were already available to us, not preregistered, and reported as such; that margin is about 3.5% of the 280 ms decision cost.

5.5. Reproducibility

Every number, table and figure in this article is generated from the released traces by `analysis/derive_results.py`, which reads the raw traces, recomputes the statistical summary and emits both a machine-readable results file and the LaTeX macros and table bodies this manuscript expands; a repository test asserts that every macro matches the corresponding derived value, so nothing measured is transcribed by hand and the guarantee is checked rather than asserted. The pipeline also emits a SHA-256 digest, byte count and row count for each of the 10 input traces (377 414 rows total), and resampling uses a fixed seed. The repository carries the historical implementation associated with the measurements and a separate corrected package described in Section 7.1; the corrected package did not generate the historical performance results. Readers should cite the Version 12 artifact path and its post-merge commit. `bash paper_iot/V13/build_v13.sh` regenerates and recompiles everything in one command, and a data dictionary (`data/raw/README.md`) marks fields withdrawn or reinterpreted in this revision. A second driver, `analysis/run_campaign.py`, repeats the throughput, latency and decision-cost campaigns against a live Fabric network in the same schema, so an independent installation's measurements can be substituted and the article regenerated against them; it does not reproduce the energy and cryptographic-timing traces, which need the INA219 bench above.

6. Results

6.1. Field-deployment latency accounting for the write path

Both transaction classes in the trace are endorsed, ordered and committed, since as Section 4.1 explains every authorization outcome writes an audit entry. The interesting question is not which class is slower but

Table 5

Write-path latency accounting at 50 concurrent clients.

Configured peers	Total latency	Recorded authorization span	Remaining latency	Reduction in remainder vs. 4 peers
4	1204.8	321.5	883.2	0.0
8	1149.5	320.4	829.1	54.1
16	1056.9	322.9	734.0	149.2
32	811.8	316.8	495.0	388.2

how much of a write is associated with the access-control layer, and the chaincode logs a timing field on every write towards answering it: `rbac_overhead_ms`, recorded alongside end-to-end latency.

6.1.1. A descriptive account, not a stage decomposition

For each configured logical peer count, write-path latency is reported as total latency, the recorded authorization-associated span, and what is left once that span is subtracted (“remaining latency”), a plain subtraction rather than a claim that the two are independently identified Fabric stages (Scope S1, Section 7).

The recorded span, averaged over the campaign, is 320 ms; we deliberately do not call this a direct measurement of permission-evaluation cost, since the timing field’s exact instrumentation boundaries cannot be reconstructed from the deployed source (S1), and it is also larger than a complete standalone `CheckAccess` transaction, 280 ms (Section 6.2), because the two counters do not share a documented boundary. The remaining latency falls from 495 ms plus 388 ms at 4 peers to 495 ms at 32 peers, a 388 ms reduction across the tested range (Table 5); nothing here identifies which of ordering, validation, commit, or the peer-count-sensitive part of endorsement execution accounts for it, and 495 ms at the largest configuration is a baseline, not an identified floor, since the remainder was still falling, with decrements per doubling growing rather than shrinking, at 32 peers.

6.1.2. The recorded span is stable in configured peer count

That the recorded span can be treated as approximately constant is tested, not assumed. The design behind that test uses 8 independent runs per configured peer count, in an order counterbalanced across eight blocks (two complete 4x4 Latin squares, one the other’s reverse-cyclic complement), so every configuration occupies every serial position exactly twice (Supplement, Table S8). We chose this design after a design review showed an earlier five-replicate schedule could not actually give every configuration every position once: four positions do not divide evenly into five replicates.

Across the four configured peer counts, the per-configuration means span 6.1 ms with no monotonic trend (1.18 ms per doubling, $p = 0.42$, $n = 32$ runs).

The extreme configurations (4 and 32 peers, 8 runs each) differ by only 4.78 ms, with a 90% interval of $[-9.77, 0.22]$ ms entirely inside the 10 ms margin (Welch $t = 1.69$, $d = 0.85$, $p = 0.114$). A joint regression on peer count and centred run sequence finds no sequence effect on the span once peer count is in the model ($p = 0.76$), which the earlier single-run, ascending-order design could not have shown. Because the span does not shrink when peers are added, its share of the write path *rises*: from 26.6% at 4 peers to 39.0% of the smaller total at 32 peers. That is the opposite of what a designer would assume from the 4-peer figure alone.

Splitting the remaining latency into ordering, validation, commit and the peer-count-sensitive part of endorsement needs a transaction timestamped at each Fabric hand-off, which we did not collect (S1); read with that boundary in mind, Table 5 answers the question a designer actually faces, how much of the budget is associated with the access-control layer, without claiming a stage attribution the measurements cannot support.

6.2. Latency by operation class

Policy evaluations (`CheckAccess`, pooled across the 7 operation types recorded) average 280 ms (SD 40) and sensor writes 1 220 ms (SD 60, P95 1 319, P99 1 359), a 940 ms gap between operation classes that differ simultaneously in payload size, world-state access and endorsement work; consensus is common to both, so this is not the cost of consensus, and should not be confused with the write-path accounting above, which is derived within rather than across classes (full table and CDF figure, Supplement Table S5 and Figure S3). The 7 operation types are pooled rather than broken out because, across the 35 role-by-operation cells in the trace, the grant rate is uniformly between 91% and 93%: a trace in which the recorded operation and role determined the policy path would instead show cells near 0% or 100%, so these labels do not separate policy-distinct executions and do not support an operation- or privilege-specific comparison. The scripted campaign in Section 6.6, where every boundary-crossing request was refused, is the evidence that the policy itself discriminates; whether privilege level or hierarchy depth affects decision *cost* needs a designed experiment (Section 7).




Figure 4: Mean throughput and standard deviation across five interleaved runs per concurrency level.

6.3. Interleaved throughput experiment

Figure 4 reports throughput against offered concurrency for the HRBAC configuration and the access-control-free baseline on identical hardware and workload.

The primary analysis is a single blocked factorial regression over every run, with terms for condition, concurrency, their interaction, and a position covariate, rather than ten independent per-level Welch tests. The ten concurrency levels were tested in one fixed ascending order, not counterbalanced across the session, so a run's raw campaign-wide sequence number is almost collinear with its concurrency level ($r = 0.995$ in this dataset) and cannot serve as an independent order covariate; only the two conditions are interleaved within each level. The position term below is therefore each run's centred position *within its own concurrency level's 10-run cell* (five interleaved HRBAC/baseline pairs). That construction makes it orthogonal to concurrency, and it is where a within-session warming or ledger-growth trend would actually show up.

The model estimates a 6.03 TPS lower throughput on the access-control-enabled path at mean concurrency (95% CI [3.20, 8.87] TPS, $p < 0.001$). The condition-by-concurrency interaction, 0.007 TPS per client, is not detectable (95% CI [-0.091, 0.104] TPS per client, $p = 0.89$), and there is no within-level position effect either ($p = 0.92$, model $R^2 = 0.206$). The model does not test for a trend across the session's fixed level order, since the design cannot separate that from the concurrency effect itself. Averaged across the ten tested concurrency levels this is a 9.6% reduction (SD 1.6 points, range 7.9–12.4%; full per-level table, Supplement Table S6). A nonsignificant interaction is not proof that the relative difference stays constant outside the tested range. It means only that no condition-by-concurrency interaction was detected within the precision of this experiment, across ten closed-loop concurrency levels from 10 to 100 clients. For capacity planning this is nonetheless the useful reading: across the levels actually tested, budgeting roughly a tenth of platform throughput for authorization held up.

Both configurations also reached their highest observed throughput at the same tested concurrency, 70.4 TPS for the baseline and 63.5 TPS for HRBAC at 40 clients, after which both decline, HRBAC by 15.3% out to 100 clients. This is the highest measured throughput under the tested closed-loop workload, not an independently established capacity ceiling; concurrency was sampled in steps of ten, so the true peak could sit anywhere in the adjacent interval, and the shared tested peak does not by itself identify which resource saturates (Scope S3, Section 7) ; ordering queueing is one candidate, reported for Fabric on x86 by Thakkar et al. [41], but so are CouchDB service times, validation and commit capacity, endorsement queueing and gateway networking, and these clients are also closed-loop, so offered concurrency is a client count rather than a controlled arrival rate. The decline past the observed peak is consistent with queueing beyond the tested workload range, but the trace does not identify the responsible resource. Admission control or added capacity are therefore engineering options to test, not conclusions established by this experiment.

6.4. Confirmatory peer-multiplicity experiment

Endorsing peers are normally added to raise physical capacity. Here we increase configured logical peer multiplicity on four fixed physical hosts (Supplement Table S3) and observe write latency, not capacity, respond. Figure 5 reports write latency as configured peer count increases from 4 to 32 at fixed offered concurrency, across 8 independent runs per configuration executed in counterbalanced order (Supplement Table S8), replacing a single-run-per-configuration design that could not separate a peer-count effect from a time-correlated one; the full latency table is Supplement Table S7.

Mean write latency falls from 1205 ms to 812 ms (run-level means, 8 runs each), a reduction of 32.6% (mean difference 393.0 ms, 95% CI [374.5, 411.4] ms), and P95 from 1297 ms to 906 ms. We report the mean difference and its interval as the primary estimate; the Welch t and Cohen's d for the same comparison




Figure 5: Run-level write latency and 95% confidence intervals by configured logical peer count.

are large for a mechanical reason (the between-run standard deviation at fixed peer count is small relative to a genuinely large peer-count effect, not transaction-level pseudo-replication: every value entering this comparison is already a run-level mean), so we report them alongside all 32 individual run-level means in Supplement Table S7 rather than in this paragraph. The decline is close to log-linear at -127.1 ms per doubling with $R^2 = 0.893$ across the four configuration means; the same effect, estimated jointly with run order on all 32 individual runs (Section 5.4), remains significant at $p < 0.001$ while run order does not ($p = 0.48$), evidence a single-run, ascending-order design could not have produced. This campaign measured latency only, not throughput, as a function of peer count (Scope S2, Section 7); the separate, interleaved concurrency-throughput comparison of Section 6.3 used a fixed peer configuration and is not informative about peer-count effects.

One hypothesis is that greater logical endorser multiplicity changed proposal routing or queueing: at four peers each zone has one endorser and concurrent requests from that zone would serialise behind it, while at 32 peers each zone has eight. Peer queue depth, proposal allocation, CPU utilisation and state-database service time were not recorded, so the experiment does not identify the mechanism behind the association. And every configured peer count above four is hosted on the same four physical machines as the four-peer configuration, so the comparison varies logical endorser multiplicity, not genuine physical distribution (S2). For irrigation control with a two-second responsiveness budget [21], greater logical peer multiplicity was associated with a larger observed margin against that threshold in this configuration; whether this association persists under other workloads, physical deployments or peer topologies remains untested.

6.5. Revocation over a multi-week window

Revocation is where short-duration benchmark measurements are least informative, because the mechanism fires rarely and its failure mode is delay

rather than error. Over the deployment, 200 lifecycle *stage records* were logged across 58 days; these are stage observations rather than independent revocations, and Section 7.5 explains why end-to-end exposure cannot be recovered from them. On the happy path the pipeline is quick: CRL publication completes in 60.4s on average (SD 4.8, maximum 71.6), and gossip to all peers takes 4.21 min (SD 0.38, P95 4.81, maximum 5.54), a tight unimodal propagation distribution with no heavy tail (Supplement, Table S9 and Figure S4).

The longitudinal view is less comfortable. 59 of the 200 stage records, 29.5%, completed in a delayed state; we give the proportion without a confidence interval, since correlated stages of the same episode are not independent trials, and CRL publication is the weakest stage at 39.1% delayed against 23.9% for gossip completion. Delayed stages are not slower on the metrics we can observe, since our timers start when publication or gossip begins and any interval spent waiting for a gateway to return is outside the measurement; taken with the connectivity outages of Section 6.8, the pattern points to scheduling and reachability rather than propagation degrading under load, though confirming it needs end-to-end timing the trace lacks (Scope S5, Section 7).

The trace contains 973 operational access attempts recorded as denied for revocation, 16.0 per day, spread out rather than concentrated in any single incident, plus a further 1340 in the scripted campaign of Section 6.6, kept separate since one is field observation and the other scripted testing. Because the deployment did not collect the denominator of all post-revocation access attempts, these counts do not estimate the probability that a revoked credential is successfully blocked; we report them as evidence enforcement fired repeatedly in the field, not as a success rate. Nor can the stage records determine, for any given delayed episode, whether disconnection was the cause, quantify end-to-end exposure from request to universal enforcement, or rule out the propagation protocol itself as a contributor: only the association with the connectivity outages of Section 6.8 is available, not a per-episode cause. For

a connected gateway, the design (Supplement, Algorithm S3) intends revocation staleness to be bounded by gossip propagation and cache invalidation, a few minutes; for an unreachable gateway, the scheduled CRL refresh interval permits a theoretical delay of up to 24 hours. The stage-level trace cannot determine how often complete revocation episodes followed either branch, or measure the resulting end-to-end exposure: only that 29.5% of individual stage records carried the deployment's delayed label: so we treat the unresolved upper bound itself, not a measured exposure rate, as the principal operational weakness the present design leaves open.

6.6. Authorization boundary enforcement

The security evaluation targets the property the system exists to provide, that no principal acts outside its role, zone or validity window. 8000 scripted attempts ran, each crossing one specific boundary of the model in Section 3, organised at collection time into 8 named scenarios of 1000 attempts. We do not report the results by those scenario names, since repetitions of a scripted request are not independent draws from a population of attacks (no interval is attached to the block rate), and, more seriously, the scenario labels do not describe the operations recorded against them; the full breakdown groups every attempt by the denial mechanism actually recorded instead (Supplement, Table S11).

All 8000 attempts were refused, invoking 6 distinct denial mechanisms rather than concentrating in one check, each spanning every role and zone in the deployment: a suite in which every attempt is caught by the same check would demonstrate one working mechanism, not a working model. So the campaign is evidence that the deployed policy denied what it was asked to deny; it is not a scenario-by-scenario conformance result, and, combined with Section 7.1, it is not evidence about the intended design either. A further scope boundary applies regardless of labelling: the expected-permission oracle against which each attempt is scored is itself derived from `roles.go`, so a defect shared between the implementation and the oracle, such as the inherited audit access of Section 7.1, would not appear as a failure here; the campaign shows that the tested requests activated 6 recorded denial mechanisms in the deployed implementation; because expected outcomes were derived from that same `roles.go`, it does not independently establish consistency with the intended policy, completeness of enforcement, or correctness for untested requests (Scope S4, Section 7). We additionally cross-check the full 8000-attempt corpus against `chaincode/policy-requirements.json`: the same independent, requirements-derived oracle built for the role hierarchy (Section 7.1). On the one dimension a static role/operation matrix can speak

to: whether the attacking role holds the requested operation as a permission at all. 2666 of 8000 attempts carry a denial reason that makes this claim (`ROLE_INSUFFICIENT` or `OP_NOT_PERMITTED`); the remainder test zone, certificate, replay or intrusion-detection logic the role/operation matrix has no opinion on. Of those 2666, 2149 (80.6%) are confirmed: the independent oracle agrees the attacking role lacks that permission. 517 (19.4%) are not: all six recorded (role, operation) pairs among them: Sensor and Gateway attempting `WriteSensor`, Farmer and Gateway attempting `ControlZone`, Certifier attempting `ReadAudit`: are pairs the independent oracle says the role *should* hold, so the recorded reason cannot be a pure role/operation question; the corpus does not carry a separate zone- or resource-ownership attribute that would let us confirm the alternative (a same-permission, wrong-zone or wrong-resource request), so we report the discrepancy rather than resolve it. This extends the independent-oracle check from the role hierarchy alone to the full boundary-attempt corpus's role/operation dimension; it remains scored against `roles.go` directly for the zone, temporal, replay and identity dimensions the corpus does not separately encode, and the corpus contains no permitted (positive) requests, only denials, so no false-denial rate can be estimated from it either. S4 is narrowed further but not closed. Transport-level attacks, physical tampering and side-channel extraction are not exercised; mutual TLS and constant-time primitives address them architecturally, not empirically.

6.7. Transmission energy and cryptographic cost

Reliability of the residue transport is withdrawn (Section 7.2); what the campaign measures is cost. Per-event transmission energy falls 48.7% [48.6, 48.8] from 24.20 mJ in single-channel mode to 12.41 mJ with residue encoding, and the signing pipeline (900 μ s) is negligible against the 30-minute duty cycle; the full breakdown, and the qualifications that bound what the energy figure describes (instrumented payload size, single-node sampling, and the small share of daily energy transmission represents), are in the supplementary material.

6.8. Availability

The online decision path, defined as reachable for *all* zones, was available 99.36% of the deployment across 3 single-gateway connectivity outages (0.2–8.7 h, 9.3 h total: a LoRa link failure, a power fluctuation and a firmware watchdog reset). The log identifies one affected gateway per event; it does not independently establish sensor-level continuity for the affected or unaffected zones. Crediting gateway-hours instead of charging the whole system for one gateway's outage gives

99.84%. Neither figure is a statement about request success rates, which we did not record per request. The sensor table contains exactly 2400 rows per day, but it has only one timestamp and no capture, gateway-receive, commit, retry or buffer-source field. It therefore cannot establish whether an outage-period row was captured before disconnection, buffered and replayed, or reconstructed during processing. The outage log's buffering notes and the regular row count are consistent with eventual replay, but they do not prove zero sensor-data loss or uninterrupted service; we report availability and ledger-row completeness separately.

7. Discussion

The implementation audit compared the released code with its specification. It identified four defects and one observability gap that the performance benchmarks did not expose. The residue transport and signature path are corrected in the released firmware and gateway code. The historical role hierarchy remains unchanged because `chaincode/hrbac/` is deliberately left exactly as deployed, and the fix is released separately (Section 7.1), together with a fourth fix, Section 7.4 below, found while preparing this revision.

7.1. Authorization-policy findings

The role hierarchy in `roles.go` places Certifier on the human inheritance chain rather than beside it, so Agronomist and Farmer receive `ReadAudit` through inheritance (Equation (1), Table 2), and Certifier itself holds `ReadZone`, so an auditor could read zone sensor data directly: the separation the introduction motivates was never in force for any of the 61 days. A second defect, found reviewing an earlier draft of the correction below, is independent of the first: Agronomist also holds `ControlZone` directly, inconsistent with the “no actuation rights” description in Section 1.

We have not patched `chaincode/hrbac/roles.go`, since doing so would make the code that produced every measurement here diverge from what is released as having produced it. Instead, `chaincode/hrbac-corrected/` fixes both. Certifier is removed from Agronomist's and Farmer's ancestor sets, and `ReadZone` moves to their direct grants. `ControlZone` is removed from Agronomist's direct grant; Farmer no longer inherits it from Agronomist, so Farmer receives its own direct grant instead. Every other effective permission is unchanged. Three regression tests verify this. `TestAuditorSeparationHolds` and `TestAgronomistCannotControlZone` each fail against `chaincode/hrbac`'s hierarchy and pass against the corrected one, confirmed with a throwaway copy run against the deployed package. `TestMatchesRequirementsOracle` checks the corrected hierarchy cell by cell, across all seven roles and ten permissions, against

Table 6

Pairwise recovery bounds for the deployed residue moduli.

Residues present (missing modulus)	Recovery bound
97, 101 (missing 103)	$97 \dot{\cap} 101 = 9\,797$
97, 103 (missing 101)	$97 \dot{\cap} 103 = 9\,991$
101, 103 (missing 97)	$101 \dot{\cap} 103 = 10\,403$

`chaincode/policy-requirements.json`, a matrix authored from the manuscript's stakeholder requirements independently of either `roles.go`, and fails and passes the same way: the independent-oracle check required by Section 6.6 (S4, Section 7). The correction is a policy patch, not a performance re-run (S5); Section 6.6 confirms only that the tested requests activated the *deployed* implementation's own recorded denial mechanisms, not that it validates the intended design.

A performance re-run against a live Fabric network is still not done, but one narrower comparison now is. A Go benchmark calls `CheckAccess` directly against the same in-memory mock state both packages' test suites already use, holding the workload and hardware fixed while only the role hierarchy differs: one grant, one deny, 2000 calls each, repeated for 5 runs (reproduce with `go test -bench=CheckAccess -benchtime=2000x -count=5` in each package; raw benchmark output is in `V13_CHANGELOG.md`, since this is a software microbenchmark, not a field or network trace, and is reported directly rather than through the derived-numbers pipeline that governs the rest of this article). Grant-path execution averages 282.9 μ s (SD 11.8) for `chaincode/hrbac` against 276.9 μ s (SD 4.9) for `chaincode/hrbac-corrected` (Welch $t = 1.06$, $p = 0.33$); deny-path execution averages 277.3 μ s (SD 5.8) against 281.9 μ s (SD 3.8) ($t = *1.48$, $p = 0.18$). Neither difference is distinguishable from run-to-run noise at this sample size. This is Go execution time inside a mock stub, not endorsement, ordering or commit against a real network, so it bounds only the chaincode's own share of the question in Section 7's scope table (S5): whatever the corrected hierarchy costs relative to the deployed one on a real network, it is not measurable in this narrow comparison.

7.2. Residue-recovery findings

As Section 4.2 sets out, recovery from an arbitrary pair of residues is unique only below the smallest pairwise product of the moduli. The deployed firmware used $\{97, 101, 103\}$, so the recovery bound depends on which residue is missing, as Table 6 shows. The trace does not identify the missing residue, so a per-pair count is unavailable.

We reconstructed the packed value for all 146400 committed readings from the released engineering-unit columns using the exact encoding in `esp32/main/main.c`: the trace's minimum is

178 382 and maximum 466 795, so 146 400 of 146 400 (100.0000%) exceed the smallest bound. 131 795 readings (90.0%) arrived with two of three residues; since every packed value here exceeded even the largest bound, each such reconstruction deterministically returned the value reduced modulo its pair's bound rather than the original, regardless of which residue was missing. All 131 795 two-residue reconstructions are therefore mathematically incorrect, and the original cannot be recovered after the fact, since raw residue packets were not retained.

The corrected firmware's $2^{253}, 2^{254}, 2^{255}$ moduli (bound 64 262) are checked by exhaustive enumeration too: every legal input to the corrected 8-bit soil encoding gives a maximum packed value of 51 455 ($2^{255} \div 201 + 200$, the full legal domain), below the bound, so every value the corrected encoder can produce is uniquely recoverable. Quantising soil moisture from twelve to eight bits to reach that bound introduces a maximum quantisation error of 0.37% of full scale (RMS 0.21%). Whether this satisfies a particular agricultural application depends on its measurement-accuracy requirement; no manufacturer tolerance claim is inferred here.

7.3. Signature-verification findings

The defect above would have been caught downstream had the signature path been correct; it was not. The firmware signs SHA-256 over a packed binary struct holding device identifier, reading identifier, raw soil moisture, temperature and timestamp, but the gateway verified an Ed25519 signature over an ASCII rendering built from the *decoded* value, a different message, so verification could not succeed against a genuine sensor signature; the same routine also returned success whenever a signature was absent, a fail-open default intended to let the gateway run without production credentials. Taken together, the 100.0% signature-validity figure in the trace cannot mean that sensor signatures were checked and passed, and we withdraw the integrity claim rather than qualify it.

The firmware signs the 32-byte SHA-256 digest of the packed struct with Ed25519 rather than using standard Ed25519ph, which uses SHA-512. The corrected verifier SHA-256-hashes the reconstructed bytes before checking the signature, matching `esp32/main/signing.c` exactly, and the test suite includes a fixed, published byte-for-byte vector an independent reimplement can check against without running this codebase (Supplement, "Signature test vector"). This is a legacy-compatible corrected verifier, not a migration to standard Ed25519 or Ed25519ph semantics: the deployed firmware's prehash construction cannot be changed without reflashing physical sensor nodes this study does not have further access to, so the correction matches what that

firmware actually signs rather than replacing it with a scheme-versioned protocol.

7.4. Certificate-revocation enforcement gap

A fourth defect surfaced while preparing this revision, not during the original audit. `CheckAccess` reads a `crl:-`prefixed world state key to decide whether a caller's certificate has been revoked (Section 3.3), but no transaction in the deployed chaincode ever wrote such a key. The operational revocation workflow (`scripts/generate-crl.sh`) already tried to call an `UpdateCRL` transaction as part of every revocation, but the deployed package has no function by that name: the call failed with an unknown-transaction error every time, the script logged the failure and moved on, and the on-ledger certificate-revocation branch of the decision algorithm was dead code for the full 61 days. This did not leave the deployment unprotected. Certificate revocation still worked, through the certificate authority's own list, enforced at the mutual-TLS layer independently of the chaincode, and through `RevokeRole` deleting the role assignment outright, a separate mechanism this article's revocation results already describe. None of the reported performance or security numbers depend on the missing transaction, because nothing in the deployed workflow ever exercised it.

`chaincode/hrbac-corrected` adds `UpdateCRL`: it parses a certificate revocation list in the format `scripts/generate-crl.sh` already produces, writes the corresponding `crl:-`prefixed state entries, and is gated to the Admin role like the package's other administrative transactions. `TestUpdateCRLWritesRevocationEntriesAndDeniesCheckAccess` and `TestUpdateCRLRequiresAdmin` in that package's test suite are the regression tests. This closes the gap for a future deployment; it cannot retroactively enforce a check the historical trace shows was never enforced.

7.5. Revocation observability

The revocation trace records lifecycle stages with no key linking them into episodes: of 132 subjects, 0 carry all four stages and 87 appear once. End-to-end exposure, the interval from revocation request to enforcement everywhere, is the quantity an operator needs, and the historical trace cannot reconstruct it (S5).

This revision adds the correlation identifier the gap called for, rather than only naming it as future work. `RevokeRole` and `UpdateCRL` in `chaincode/hrbac-corrected` both stamp an `EpisodeID` field on their audit entries, taken from the same nonce value `scripts/revoke-cert.sh` and `scripts/generate-crl.sh` now share for one revocation request. On the gateway side, `PolicyCache.invalidate/invalidate_all` log that same token when a cache is cleared, and a new `Gateway.handle_revocation` method gives

an event-driven listener a real, tested place to call. Before this change, nothing in the repository ever called `PolicyCache.invalidate`, so the event-driven invalidation this article describes had no caller anywhere in this codebase; a live deployment's gossip or chaincode-event subscription would still need to be wired to that method, which is outside what this revision changes. What this gives a future deployment is a real, joinable key across the chaincode and gateway sides of one revocation. It does not, and cannot, resolve episodes already present in the historical trace analysed here: that data was collected before the correlation identifier existed.

7.6. Cross-layer implications

The defects occurred at interfaces between components: a role hierarchy disagreeing with its own description, a transport whose recovery bound the encoder never checked, a signature verified over the wrong bytes, and a ledger check with no writer. The artifact now generates the permission table from `roles.go` instead of transcribing it, checks every legal encoder input against the recovery bound, tests gateway verification against the firmware's byte layout, and tests that `UpdateCRL` actually denies a caller whose certificate it revokes.

Table 7 states, for each finding, what a reader may still treat as valid and what they may not.

7.7. Performance implications

Four consequences follow, each contradicting a plausible default assumption.

The recorded authorization-associated span accounted for roughly a quarter of the write path, and it varied less across tested peer configurations than the uninstrumented remainder did. It sits at 26.6% of the write path at 4 peers and does not move appreciably as configured peer count scales, so optimising policy evaluation is bounded above by roughly a quarter of write latency (Scope S1). Because expired-role cleanup mutates state inside `CheckAccess`, read-intent checks are endorsed rather than served as peer queries. A read-only decision path, with expiry reclaimed by a background sweep instead, would take reads off the consensus path, at the cost of the durable ordered audit transaction a committed `CheckAccess` currently produces, and it would need revocation freshness maintained another way. We call that a high-value candidate worth the trade-off in most deployments, not an unconditional recommendation.

Peers are associated with lower tail latency here; treat throughput separately. Configured logical peer multiplicity is associated with a close-to-log-linear reduction in write latency (Section 6.4), supported by a joint regression on peer count and run order once the earlier order confound is designed out (S2). The separate, interleaved concurrency experiment, meanwhile, found the same observed throughput peak for both

Table 7

Validity and permitted interpretation of evidence affected by the implementation audit.

Evidence	Status	Defensible interpretation
Authorization decisions, denial mechanisms (Section 6.6)	Valid for the deployed policy	A record of what the <i>deployed</i> hierarchy enforced, not the intended one
Sensor readings recovered from two of three residues (90.0%)	Withdrawn	Mathematically incorrect values; not recoverable after the fact
Signature-validity field	Withdrawn	Cannot be read as "sensor signatures were checked and passed"
On-ledger certificate-revocation check (Section 7.4)	Never enforced	<code>CheckAccess</code> 's CRL branch had no writer; revocation still worked via the CA's list and role deletion, which the reported results reflect
Revocation end-to-end exposure	Not measurable from this trace	Only per-stage durations are usable; a correlation identifier now exists in chaincode/ <code>hrbac-corrected</code> (Section 7.5) for future collection, but was not present when this trace was recorded

arms without locating the saturating resource (S3). The phrase "peers buy latency, not capacity" describes only the latency side, and a deployment weighing more endorsers against more capacity should measure both under its own workload.

Budget the access-control-enabled path empirically, not as a universal constant. Across the tested closed-loop concurrency levels, the complete access-control-enabled path averaged 9.6% lower throughput than the comparison path, and no condition-by-concurrency interaction was detected. That supports using roughly that figure as a workload-specific planning estimate for the tested peer count and workload, but it does not establish a fixed derating factor that transfers to other peer counts, workloads or high-load regimes.

Revocation is the weak link. It is the only mechanism whose behaviour materially degraded over the deployment: 59 of 200 stage records carry the deployment's delayed label. Their association with recorded connectivity outages motivates a reachability hypothesis, but the trace cannot determine per-episode cause or exclude the propagation protocol as a contributor

(S5). Event-driven credential status checking [47] and constrained-device status protocols such as TinyOCSP [8] remove the dependency on a scheduled publication reaching a possibly-disconnected gateway, whichever the cause turns out to be; this revision's episode-correlation identifier (Section 7.5) is a smaller step in the same direction, aimed at measuring the problem rather than fixing it outright.

Direct numerical comparison against prior systems is limited by methodological heterogeneity, hardware, workload, transaction semantics and what is counted as latency all differ (Supplement, Table S1), a point Abang et al. [1] make about the Fabric latency literature generally; the comparison we regard as meaningful is the internal one, the same network, hardware and workload with and without the access-control layer (Section 6.3), since it is reproducible from the released artifact.

7.8. Scope of validation

The threat model treats orderers and the certificate authority as non-malicious under crash-fault-tolerant Raft, and physical tampering and side-channel extraction as residual risk, architecturally addressed by mutual TLS and constant-time primitives rather than empirically tested. Within that scope, three things this deployment measured are not the same as validating the intended design. Section 6.6 measures refusals against the *deployed* policy, which Section 7.1 shows is not the intended one, and its oracle is derived from the same code it tests (S4), so it is conformance testing of one implementation, not a security proof, and says nothing about an adversary who deviates from the scripted campaign. The audit's corrections are patches to code, not to the archived timing measurements. A chaincode-only microbenchmark (Section 7.1) finds `CheckAccess` execution cost statistically indistinguishable between the deployed and corrected role hierarchies, but that isolates Go execution time inside the `mockStub` harness used to test both packages; it says nothing about whether the corrected chaincode preserves the write latency, throughput and recorded span this article reports for a real Fabric network, which has not been re-run (S5). And overhead throughout is attributed by differencing against a baseline with the access-control chaincode removed, which also removes the audit write and the nonce check, so every reported cost bounds the access-control *layer* rather than an isolated permission decision. Table 7 in Section 7 states, trace by trace, what remains usable after the audit.

7.9. Scope and validity

Table 8 defines the evidence boundaries and the measurements needed to extend the study scope.

One deployment, one season, one site, and a workload dominated by a regular sensor-write stream: we regard the recorded span's small, stable share of the write path as more transferable than the millisecond

values (S2). The topology trace's coordinate columns are unreliable and unused, and a multi-device energy campaign over the full authenticated frame would generalise Section 6.7's single-node figure.

Table 8

Evidence boundaries and measurements required to extend the study scope.

	Scope boundary	Extension required
S1	The recorded span's instrumentation boundaries are undocumented, so it is not comparable to standalone CheckAccess latency, and the remaining latency stays aggregated over ordering, validation, commit and peer-insensitive endorsement work; operation and role labels in the trace also do not separate policy-distinct executions, so whether privilege level costs anything is untested.	Both counters specified at the instruction level; a transaction timestamped at each Fabric hand-off; a designed experiment varying hierarchy depth and denial reason with workload fixed.
S2	The counterbalanced peer-scaling campaign controls for run order within one follow-up campaign on the same four physical hosts used throughout; it does not establish the result on different hardware or in a different season, and it measured latency only, not throughput, as a function of peer count.	Configurations on physically distinct hosts, across seasons, with throughput measured alongside latency.
S3	The interleaved concurrency benchmark used a fixed peer configuration and ten-client sampling steps, so it cannot locate the saturating resource or the true peak between tested levels.	Orderer queue depth, block-fill rate and CouchDB service times under a controlled intervention.
S4	The 8 000-attempt authorization-boundary campaign is scored against a permission oracle derived from the same roles.go it tests, so a shared defect would not surface there; it is conformance testing of the deployed policy, not validation of the intended design or resistance to an adaptive adversary. An independent, requirements-derived oracle (policy-requirements.json, Section 7.1) now cross-checks the full corpus on the role/operation dimension (2666 of 8 000 attempts; 80.6% confirmed, 19.4% unresolved because the corpus lacks a zone/resource attribute), but the corpus contains no permitted requests and the zone, temporal, replay and identity dimensions remain scored against roles.go	A permitted-request (not only denial) corpus; per-attempt zone and resource-ownership attributes so the unresolved cases can be adjudicated; mutation tests on inheritance edges; randomised testing against a defined attack population.

8. Conclusion

This study evaluated blockchain access control during a 61-day agricultural edge-IoT deployment and two controlled campaigns. The recorded authorization-associated span averaged 320 ms within a 1 205 ms write path. Increasing configured logical peers from 4 to 32 on four fixed hosts reduced the uninstrumented remainder by 388 ms. The interleaved throughput experiment found that the access-control-enabled path averaged 9.6% lower throughput across the tested range. These findings apply to the evaluated topology and workload.

The implementation audit identified policy, residue-recovery, signature-verification and certificate-revocation defects, together with a revocation-observability gap this revision partly closes in code with a shared correlation identifier. The released corrections preserve the historical implementation for reproducibility. A chaincode-only microbenchmark finds no measurable decision-cost difference between the deployed and corrected role hierarchies, but network-level performance under the corrected chaincode has not been re-benchmarked. The public artifact provides the traces, schedules, code, tests, and analysis needed to reproduce the results and evaluate the design in other settings.

CRedit authorship contribution statement

Forcha Glen Beloa: Conceptualization, Methodology, Software, Investigation, Data curation, Formal analysis, Writing: original draft. **Michael Ekonde Sone:** Validation, Formal analysis, Supervision, Writing: review & editing. **Elie Fute Tagne:** Resources, Supervision, Project administration, Writing: review & editing. **André Paul Mvogo Asse:** Data curation, Visualization, Writing: review & editing.

Ethics and data protection

The traces record authorization decisions whose subjects include farm operators and administrators, so caller identifiers are potentially personal data. Oversight for the deployment was provided by the University of Buea Faculty of Engineering and Technology's departmental research ethics review, which issued a non-interventional technical study determination (reference FET-CE-2025-014, 14 July 2025) covering the collection of operational, performance and access-control traces from consenting farm personnel; no biomedical or behavioural data were in scope. Each participating operator and administrator gave written informed consent before the deployment began, covering data collection, the use of pseudonymised traces in this article and their public release; participation was voluntary and could be withdrawn without affecting farm operations.

Identifiers in the released dataset are pseudonyms assigned at collection (`usr-` and `sns-` prefixes); no names, contact details or free-text fields were collected, and the mapping from pseudonym to individual was destroyed by the data-collection team before the dataset was prepared for release, so no reverse mapping exists for us or for a reader to reconstruct. The site is identified at settlement granularity rather than by field coordinates, and the coordinate columns in the topology trace are unreliable in any case (Section 7.9). No sensor data revealing an individual's conduct beyond routine farm operation was collected. These are the grounds on which the pseudonymised operational traces are released alongside this article.

Funding

This work received mobility support from the MIRET (Mobility for Innovation and Research in Emerging Technologies) Intra-Africa Programme. That support funded travel and exchange activity only; it had no role in study design, data collection, analysis, interpretation, manuscript preparation or the decision to submit. No other specific grant was received from funding agencies in the public, commercial or not-for-profit sectors.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Declaration of generative AI in the writing process

During the preparation of this work the authors used an AI coding assistant for language editing and for drafting code in the analysis pipeline (`analysis/derive_results.py`) and its statistical routines. No AI-generated code contributes to a reported result without independent verification: every statistical routine is exercised by the test suite in `analysis/test_manuscript_consistency.py`, which recomputes every macro, table and figure from the raw traces and fails on any mismatch; select statistics (the peer-scaling Welch comparison of Section 6.4 and the throughput factorial model of Section 6.3) were additionally reproduced with an independent, hand-written script (`analysis/independent_check.py`) using a different estimation path than `derive_results.py`'s. The authors reviewed and edited all content, including all AI-drafted code, and take full responsibility for the content of the published article and the correctness of its computations.

Acknowledgements

We thank the participating farm operators in Sfax for their collaboration and practical insight, and acknowledge support from the MIRET (Mobility for Innovation and Research in Emerging Technologies) Intra-Africa Programme.

Data availability

The trace dataset, firmware, formal specification, benchmark harness, requirements-derived security oracle and the analysis pipeline that generates every number, table and figure in this article are publicly available in the Version 13 artifact at https://github.com/forchag/project-fet2/tree/main/paper_1ot/V13. The repository preserves the historical implementation and the corrected artifacts as separate packages. Historical performance results were generated from the deployed implementation. Corrected-implementation performance on a real Fabric network was not re-benchmarked; a chaincode-only microbenchmark comparing decision cost in isolation is reported in Section 7.1. The exact Git commit containing this Version 13 artifact will be recorded in the camera-ready submission, and an archival DOI will be added only after the repository release has been minted. Until then, the repository path above is the citable public source.

References

- [1] Abang, J.E., Takruri, H., Al-Zaidi, R., Al-Khalidi, M., 2024. Latency performance modelling in hyperledger fabric blockchain: Challenges and directions with an IoT perspective. *Internet of Things* 26, 101217. doi:10.1016/j.iot.2024.101217.
- [2] Adelantado, F., Vilajosana, X., Tuset-Peiro, P., Martinez, B., Melia-Segui, J., Watteyne, T., 2017. Understanding the limits of LoRaWAN. *IEEE Communications Magazine* 55, 34–40. doi:10.1109/MCOM.2017.1600613.
- [3] Alshehri, A., Sandhu, R., 2016. Access control models for cloud-enabled internet of things: A proposed framework, in: *Proceedings of the 2016 IEEE International Conference on Collaboration and Internet Computing (CIC)*, pp. 530–538. doi:10.1109/CIC.2016.075.
- [4] Androulaki, E., Barger, A., Bortnikov, V., Cachin, C., Christidis, K., De Caro, A., Enyeart, D., Ferris, C., Laventman, G., Manevich, A., Muralidharan, S., Murthy, C., Nguyen, B., Sethi, M., Singh, G., Smith, K., Sorniotti, A., Stathakopoulou, C., Vukolić, M., Cocco, S.W., Yellick, J., 2018. Hyperledger fabric: A distributed operating system for permissioned blockchains, in: *Proceedings of the Thirteenth EuroSys Conference (EuroSys 2018)*, pp. 30:1–30:15. doi:10.1145/3190508.3190538.
- [5] Bandara, E., Liang, X., De Zoysa, K., Foytik, P., Shetty, S., 2021. Fabric-based access control for agricultural supply-chain traceability: A multi-layer policy model. *IEEE Transactions on Dependable and Secure Computing* 18, 2089–2101. doi:10.1109/TDSC.2021.3092089.
- [6] Benedict, S., 2021. Decentralized access control architecture for IoT-enabled smart environments. *Internet of Things* 14, 100379. doi:10.1016/j.iot.2021.100379.
- [7] Castro, M., Liskov, B., 1999. Practical byzantine fault tolerance, in: *Proceedings of the Third Symposium on Operating Systems Design and Implementation (OSDI 99)*, pp. 173–186.
- [8] Fedrecheski, G., Rabaey, J.M., Costa, L.C.P., Calcina-Ccori, P.C., Pereira, W.T., Zuffo, M.K., 2023. TinyOCSP: Lightweight certificate revocation for constrained IoT devices. *Internet of Things* 22, 100740. doi:10.1016/j.iot.2023.100740.
- [9] Ferdous, M.S., Chowdhury, F., Allassafi, M.O., 2019. In search of self-sovereign identity leveraging blockchain technology. *IEEE Access* 7, 103059–103079. doi:10.1109/ACCESS.2019.2931173.
- [10] Ferrag, M.A., Shu, L., Yang, X., Derhab, A., Maglaras, L., 2020. Security and privacy for green IoT-based agriculture: Review, blockchain solutions, and challenges. *IEEE Access* 8, 32031–32053. doi:10.1109/ACCESS.2020.2975715.
- [11] Ferraiolo, D.F., Sandhu, R., Gavrila, S., Kuhn, D.R., Chandramouli, R., 2001. Proposed NIST standard for role-based access control. *ACM Transactions on Information and System Security* 4, 224–274. doi:10.1145/501978.501980.
- [12] Garner, H.L., 1959. The residue number system. *IRE Transactions on Electronic Computers* EC-8, 140–147. doi:10.1109/TEC.1959.5219515.
- [13] González, V., Sánchez, L., Lanza, J., Santana, J.R., Sotres, P., García, A.E., 2023. On the use of blockchain to enable a highly scalable Internet of Things data marketplace. *Internet of Things* 22, 100722. doi:10.1016/j.iot.2023.100722.
- [14] Hernández-Ramos, J.L., Jara, A.J., Marín, L., Skarmeta, A.F., 2013. Distributed capability-based access control for the internet of things. *Journal of Internet Services and Information Security* 3, 1–16.
- [15] Hu, V.C., Ferraiolo, D., Kuhn, R., Schnitzer, A., Sandlin, K., Miller, R., Scarfone, K., 2014. Guide to Attribute Based Access Control (ABAC) Definition and Considerations. Technical Report NIST SP 800-162. National Institute of Standards and Technology. doi:10.6028/NIST.SP.800-162.
- [16] Hyperledger Foundation, 2025. Hyperledger fabric v3.0: What's new. Available at <https://hyperledger-fabric.readthedocs.io/en/release-3.0/whatsnew.html>.
- [17] Idrissi, M., Palmieri, F., 2024. Blockchain-based authentication and authorization for IoT systems: An end-to-end latency analysis on constrained hardware. *The Journal of Supercomputing* 80, 1–32. doi:10.1007/s11227-024-06001-7.
- [18] Kamilaris, A., Fonts, A., Prenafeta-Boldú, F.X., 2019. The rise of blockchain technology in agriculture and food supply chains. *Trends in Food Science & Technology* 91, 640–652. doi:10.1016/j.tifs.2019.07.034.
- [19] Ke, C.K., Park, S., 2023. Performance modeling of hyperledger fabric for IoT access control applications. *Future Internet* 15, 115. doi:10.3390/fi15030115.
- [20] Lamport, L., 2002. *Specifying Systems: The TLA+ Language and Tools for Hardware and Software Engineers*. Addison-Wesley.
- [21] Lea-Cox, J.D., Belayneh, A., van Iersel, M.W., Lichtenberg, E., Kantor, G.F., 2013. Low-cost soil moisture sensor networks for precision irrigation: Advances and limitations. *HortTechnology* 23, 792–798. doi:10.21273/HORTTECH.23.6.792.
- [22] Lin, Q., Wang, H., Pei, X., Wang, J., 2019. Food safety traceability system based on blockchain and EPCIS. *IEEE Access* 7, 20698–20707. doi:10.1109/ACCESS.2019.2897792.
- [23] Liu, Y., Yu, F.R., Li, X., Ji, H., Leung, V.C.M., 2019. Distributed resource allocation in blockchain-based video streaming systems with mobile edge computing. *IEEE Transactions on Wireless Communications* 18, 695–708. doi:10.1109/TWC.2018.2886345.
- [24] Mohanta, B.K., Awad, A.I., Elsaka, T., Kheddar, H., Baraka, E., 2025. Smart-contract-based blockchain-enabled decentralized scheme for improving smart-grid security. *Internet of Things* 34, 101811. doi:10.1016/j.iot.2025.101811.
- [25] Neumann, A., Aichele, C., Lindner, M., Wunderlich, S., 2008. Better Approach To Mobile Ad-hoc Networking (B.A.T.M.A.N.). Internet Draft draft-openmesh-b-a-t-m-a-n-00. IETF. Available at <https://tools.ietf.org/html/draft-openmesh-b-a-t-m-a-n-00>.
- [26] Newcombe, C., Rath, T., Zhang, F., Munteanu, B., Brooker, M., Deardeuff, M., 2015. How Amazon Web Services uses formal methods. *Communications of the ACM* 58, 66–73. doi:10.1145/2699417.
- [27] Novo, O., 2018. Blockchain meets IoT: An architecture for scalable access management in IoT. *IEEE Internet of Things Journal* 5, 1184–1195. doi:10.1109/JIOT.2018.2812239.
- [28] Ongaro, D., Ousterhout, J., 2014. In search of an understandable consensus algorithm, in: *Proceedings of the 2014 USENIX Annual Technical Conference (USENIX ATC 14)*, pp. 305–319.
- [29] Ouaddah, A., Abou Elkalam, A., Ait Ouahman, A., 2016. FairAccess: A new blockchain-based access control framework for the Internet of Things. *Security and Communication Networks* 9, 5943–5964. doi:10.1002/sec.1748.
- [30] Pajoo, H.H., Hammoudeh, M., Alam, F., Brocklehurst, S., 2021. Hyperledger fabric blockchain for securing the edge IoT: Systems architecture, experiments and critical observations. *Electronics* 10, 1338. doi:10.3390/electronics10111338.
- [31] Putra, G.D., Dedeoglu, V., Kanhere, S.S., Jurdak, R., 2021. Trust management in decentralized IoT access control system, in: *Proc. IEEE BCCA*, pp. 1–8. doi:10.1109/BCCA.53669.2021.9657049.

- [32] Raza, S., Shafagh, H., Hewage, K., Hummen, R., Voigt, T., 2013. Lithe: Lightweight secure CoAP for the Internet of Things. *IEEE Sensors Journal* 13, 3711–3720. doi:10.1109/JSEN.2013.2277656.
- [33] Rustemi, A., Dalipi, F., 2025. Synergizing IoT, AI, and blockchain for smart agriculture: Challenges, opportunities, and future directions. *Internet of Things* 34, 101778. doi:10.1016/j.iot.2025.101778.
- [34] Safeer, S., Gallo, P., Pulvento, C., 2025. Agri-farming with computer vision, IoT and blockchain towards climate smart cultivation. *Internet of Things* 34, 101749. doi:10.1016/j.iot.2025.101749.
- [35] Sandhu, R., Bhamidipati, V., Munawer, Q., 1999. The ARBAC97 model for role-based administration of roles. *ACM Transactions on Information and System Security* 2, 105–135. doi:10.1145/300830.300839.
- [36] Sandhu, R.S., Coyne, E.J., Feinstein, H.L., Youman, C.E., 1996. Role-based access control models. *Computer* 29, 38–47. doi:10.1109/2.485845.
- [37] Sciancalepore, S., Piro, G., Boggia, G., Bianchi, G., 2017. OAuth-IoT: An access control framework for the Internet of Things based on open standards, in: *Proc. IEEE ISCC*, pp. 676–681. doi:10.1109/ISCC.2017.8024606.
- [38] Servos, D., Osborn, S.L., 2017. Current research and open problems in attribute-based access control. *ACM Computing Surveys* 49, 65:1–65:45. doi:10.1145/2983429.
- [39] Shih, D.H., Jheng, P.Y., Yen, D.C., Huang, C.L., 2022. A blockchain-based access control model for industrial IoT. *Sensors* 22, 2188. doi:10.3390/s22062188.
- [40] Tara, S., Zeadally, S., Sookun, D., 2023. Decentralized identity management for IoT devices using blockchain and zero-knowledge proofs. *IEEE Internet of Things Journal* 10, 3817–3827. doi:10.1109/JIOT.2022.3211254.
- [41] Thakkar, P., Nathan, S., Viswanathan, B., 2018. Performance benchmarking and optimizing hyperledger fabric blockchain platform, in: *Proceedings of the 26th IEEE International Symposium on Modeling, Analysis, and Simulation of Computer and Telecommunication Systems (MASCOTS)*, pp. 264–276. doi:10.1109/MASCOTS.2018.00060.
- [42] Tian, F., 2017. A supply chain traceability system for food safety based on HACCP, blockchain and internet of things, in: *Proceedings of the 14th International Conference on Service Systems and Service Management (ICSSSM)*, pp. 1–6. doi:10.1109/ICSSSM.2017.7996119.
- [43] Tzenetopoulos, A., Katsaros, P., Mavridis, I., 2022. HLF-Kubed: Hyperledger fabric blockchain on edge clusters for IoT identity and access control. *Future Generation Computer Systems* 135, 271–283. doi:10.1016/j.future.2022.05.003.
- [44] Usman, M., Jolfaei, A., Tariq, M.U., 2024. Blockchain-based access control for industrial IoT: Domain policy architecture with smart contracts. *Computer Networks* 243, 110286. doi:10.1016/j.comnet.2024.110286.
- [45] Wang, X., Han, Y., Leung, V.C.M., Niyato, D., Yan, X., Chen, X., 2020. Convergence of edge computing and deep learning: A comprehensive survey. *IEEE Communications Surveys & Tutorials* 22, 869–904. doi:10.1109/COMST.2020.2970550.
- [46] World Wide Web Consortium, 2022. Decentralized identifiers (DIDs) v1.0. W3C Recommendation, 19 July 2022. Available at <https://www.w3.org/TR/did-core/>.
- [47] World Wide Web Consortium, 2024. Bitstring status list v1.0. W3C Candidate Recommendation Snapshot, 5 March 2024. Available at <https://www.w3.org/TR/vc-bitstring-status-list/>.
- [48] Xu, R., Chen, Y., Blasch, E., Chen, G., 2018. BlendCAC: A blockchain-enabled decentralized capability-based access control for IoTs, in: *Proceedings of the 2018 IEEE International Conference on Internet of Things (iThings)*, pp. 1027–1034. doi:10.1109/Cybermatics_2018.2018.00191.
- [49] Ye, N., Zhu, Y., Wang, R.c., Malekian, R., Lin, Q.m., 2014. An efficient authentication and access control scheme for perception layer of internet of things. *Applied Mathematics & Information Sciences* 8, 1617–1624. doi:10.12785/amis/080416.
- [50] Yin, M., Malkhi, D., Reiter, M.K., Gueta, G.G., Abraham, I., 2019. HotStuff: BFT consensus with linearity and responsiveness, in: *Proceedings of the 2019 ACM Symposium on Principles of Distributed Computing (PODC)*, pp. 347–356. doi:10.1145/3293611.3331591.
- [51] Zaidi, S.M.S., Al-Qurishi, M., Abidin, M.Z.u., Almakdi, S., Asiri, Y., 2023. Fabrication: A hyperledger fabric framework for scalable IoT access control. *IEEE Internet of Things Journal* 10, 12636–12649. doi:10.1109/JIOT.2023.3254978.
- [52] Zhang, N., Wang, M., Wang, N., 2002. Precision agriculture: A worldwide overview. *Computers and Electronics in Agriculture* 36, 113–132. doi:10.1016/S0168-1699(02)00056-X.