

**UNIVERSITY OF BUEA
FACULTY OF ENGINEERING AND TECHNOLOGY
DEPARTMENT OF COMPUTER ENGINEERING**

**A DISTRIBUTED BLOCKCHAIN-BASED APPROACH TO
PROTECT IDENTITY, INTEGRITY AND PRIVACY IN
AGRICULTURAL DATA PROCESSING**

**BY
FORCHA GLEN BELOA**

(FE22P049)

M.Eng. in Telecommunications and Networks

A thesis submitted to the Faculty of Engineering and Technology of the
University of Buea in partial fulfilment of the requirements for the
degree of Doctor of Philosophy (PhD) in Computer Engineering

Supervisor

Michael Ekonde SONE, PhD
Associate Professor

Co-Supervisor

Elie FUTE TAGNE, PhD
Associate Professor

August 2026

UNIVERSITY OF BUEA
FACULTY OF ENGINEERING AND TECHNOLOGY
DEPARTMENT OF COMPUTER ENGINEERING

**A DISTRIBUTED BLOCKCHAIN-BASED APPROACH TO
PROTECT IDENTITY, INTEGRITY AND PRIVACY IN
AGRICULTURAL DATA PROCESSING**

BY
FORCHA GLEN BELOA

(FE22P049)

M.Eng. in Telecommunications and Networks

A thesis submitted to the Faculty of Engineering and Technology of the
University of Buea in partial fulfilment of the requirements for the
degree of Doctor of Philosophy (PhD) in Computer Engineering

Supervisor

Michael Ekonde SONE, PhD
Associate Professor

Co-Supervisor

Elie FUTE TAGNE, PhD
Associate Professor

August 2026

DEDICATION

To . . .

CERTIFICATION

UNIVERSITY OF BUEA
FACULTY OF ENGINEERING AND TECHNOLOGY

This is to certify that the work described in the thesis entitled **“A Distributed Blockchain-Based Approach to Protect Identity, Integrity and Privacy in Agricultural Data Processing”**

By _____

Was carried out in the Department of Computer Engineering

Under the supervision of

(Name and Signature of Supervisor)

Date _____

ACKNOWLEDGEMENTS

ABSTRACT

TABLE OF CONTENTS

.....	x
1 Introduction	1
1.1 Background to the Study	1
1.2 Statement of the Problem	3
1.3 Rationale of the Study	5
1.4 Objectives of the Study	6
1.4.1 Main Objective	6
1.4.2 Specific Objectives	6
1.5 Research Questions	7
1.6 Research Hypotheses	8
1.7 Scope and Delimitations	8
1.8 Significance of the Study	9
1.9 Definition of Key Terms	10
1.10 Organisation of the Thesis	10

LIST OF FIGURES

1.1	Lifecycle of an agricultural measurement. A reading becomes a record that is consumed by parties outside the farm, and the value those parties attach to it is what makes falsification worthwhile.	2
1.2	The three protection properties and the shortfall in each. Because all three draw on a single constrained budget, they form one design problem rather than three independent ones.	5
1.3	Organisation of the thesis and the chapter in which each specific objective is discharged.	10

LIST OF TABLES

LIST OF ABBREVIATIONS

ABAC	Attribute-Based Access Control
ACK	Acknowledgement
ADC	Analogue-to-Digital Converter
API	Application Programming Interface
ARM64	64-bit Advanced RISC Machine architecture
BATMAN-adv	Better Approach To Mobile Ad-hoc Networking, advanced
BLE	Bluetooth Low Energy
BFT	Byzantine Fault Tolerance
CA	Certificate Authority
CAP	Consistency, Availability, Partition tolerance
CPU	Central Processing Unit
CRT	Chinese Remainder Theorem
DID	Decentralised Identifier
DLT	Distributed Ledger Technology
DP	Differential Privacy
EDP	Energy-Delay Product
ESP32	Espressif 32-bit microcontroller with integrated Wi-Fi and Bluetooth
GDPR	General Data Protection Regulation
GW	Gateway
HLF	Hyperledger Fabric
HRBAC	Hierarchical Role-Based Access Control
IoT	Internet of Things
IPFS	InterPlanetary File System
JSON	JavaScript Object Notation
LoRa	Long Range (sub-GHz spread-spectrum radio modulation)
LoRaWAN	Long Range Wide Area Network
LZW	Lempel-Ziv-Welch compression
MAC	Medium Access Control
MCU	Microcontroller Unit
MSP	Membership Service Provider
PBFT	Practical Byzantine Fault Tolerance
PKI	Public Key Infrastructure
RAM	Random Access Memory
RBAC	Role-Based Access Control
RF	Radio Frequency
ROI	Return on Investment
RPi	Raspberry Pi
RTC	Real-Time Clock
SF	Spreading Factor
SHA	Secure Hash Algorithm

SLO	Service Level Objective
SSS	Shamir Secret Sharing
TEE	Trusted Execution Environment
TPS	Transactions Per Second
TPM	Trusted Platform Module
ZKP	Zero-Knowledge Proof

CHAPTER ONE


INTRODUCTION

1.1 Background to the Study

Agriculture is undergoing a transition from experience-driven practice to data-driven decision making. Networks of low-cost sensors now measure soil moisture, canopy temperature, humidity and nutrient status continuously, and the resulting streams feed irrigation controllers, yield models and certification schemes [1, 2]. In water-stressed regions this shift is not a convenience but a necessity. Irrigation accounts for the dominant share of freshwater withdrawal, and climate variability has compressed the margin for error in scheduling decisions [3, 4].

A measurement taken in a field does not stop at the field. It becomes a record, and that record is consumed by parties with interests of their own, as Figure 1.1 sets out. The value of the record to those parties is exactly what creates the incentive to falsify it, and the further a record travels from the instrument that produced it, the weaker the evidence that it is genuine.

The sensing layer of such systems is built from severely constrained devices. An Espressif 32-bit microcontroller with integrated Wi-Fi and Bluetooth (ESP32) node operating on battery and harvested solar power must transmit over kilometre-scale distances using Long Range (sub-GHz spread-spectrum radio modulation) (LoRa), a modulation scheme that trades data rate for range and therefore charges a high airtime cost per packet [5, 6]. Airtime is the binding constraint. It determines energy consumption at



*Each downstream use attaches value to the record,
and therefore creates an incentive to falsify it.*

Figure 1.1. *Lifecycle of an agricultural measurement. A reading becomes a record that is consumed by parties outside the farm, and the value those parties attach to it is what makes falsification worthwhile.*

the node, and, through duty-cycle regulation and collision probability, it determines how many nodes a single gateway can serve [7]. Conventional responses compress the payload, but general-purpose compressors such as Huffman and Lempel-Ziv-Welch compression (LZW) carry dictionary and buffer overheads that a microcontroller with a few kilobytes of free Random Access Memory (RAM) can ill afford [8, 9].

The Chinese Remainder Theorem (CRT) offers a different route. Decomposing a reading into residues modulo pairwise coprime moduli yields several short packets in place of one long one [10, 11]. Because a multi-channel concentrator such as the SX1302 demodulates channels concurrently, the short residue packets can be received simultaneously, and the original value is recovered in $O(k)$ operations for k moduli. The practical effect is a reduction in per-packet occupancy rather than in total transmitted volume, which is precisely the quantity that governs gateway capacity. The detailed encoding and reconstruction path is presented in Chapter Three.

Data that has been collected cheaply must then be trusted. Agricultural records increasingly carry consequences beyond the farm gate. They substantiate organic and fair-trade certification, they underwrite index-based crop insurance, and they inform subsidy allocation. Each of these uses creates an incentive to falsify. A centralised database

concentrates that risk in a single administrator, and an unprotected sensor network invites the injection of fabricated readings by any party with a compatible radio. Distributed Ledger Technology (DLT) addresses the tamper-evidence requirement by binding records into a hash-linked chain replicated across mutually distrusting parties, and permissioned platforms such as Hyperledger Fabric (HLF) add an identity layer through a Membership Service Provider (MSP) and Public Key Infrastructure (PKI), making it feasible to say not merely that a record is unaltered but who produced it [12–14].

Running that machinery at the edge is the difficulty. Consensus, endorsement and world-state maintenance were designed for datacentre hardware, whereas an agricultural deployment must place gateways in the field, on 64-bit Advanced RISC Machine architecture (ARM64) single-board computers, behind intermittent connectivity. Whether a permissioned ledger can meet irrigation-control timing requirements on such hardware has until recently been asserted rather than measured.

Finally, protection is not exhausted by integrity. A soil-moisture trace is also a disclosure. It reveals plot productivity, irrigation strategy and, in aggregate, the commercial position of the farmer. Where records are replicated across a consortium that includes buyers, insurers and cooperatives, the same replication that secures integrity enlarges the audience for private information. Identity, integrity and privacy therefore have to be treated as one design problem rather than three independent ones, which is the position this thesis takes.

1.2 Statement of the Problem

Agricultural data processing systems are currently forced into a choice between three properties that ought to hold simultaneously.

Identity is weak at the sensing layer. Constrained nodes seldom carry credentials

strong enough to authenticate individual readings. Where authentication exists it typically terminates at the gateway, so the ledger attests to what the gateway submitted rather than to what the sensor measured. The provenance chain has a gap exactly where falsification is cheapest, at the final hop between an individual node and the record that claims to represent it.

Integrity guarantees do not survive the move to the edge. Permissioned ledgers provide tamper evidence, but their published performance characterisations assume server-class hosts. On ARM64 gateways with constrained memory and intermittent links, it is not established that endorsement and ordering can complete inside the control loop of an irrigation decision, nor how the ledger behaves during network partition and recovery.

Privacy is sacrificed to obtain integrity. Replicating raw agronomic readings to every consortium member is the simplest way to make them verifiable and the worst way to keep them confidential. Existing deployments largely accept this trade, or push confidential data off chain in a way that reintroduces the trusted intermediary the ledger was meant to remove.

Figure 1.2 places the three shortfalls together. They are not independent faults to be repaired one at a time; each is a consequence of the same scarcity, and a remedy applied to one is paid for out of the budget available to the others.

Underlying all three is a resource problem. The transmission budget that would pay for per-reading signatures, encryption and privacy-preserving encodings is already consumed by the raw payload itself. Any scheme that adds cryptographic protection without first reducing transmission cost will exhaust the node's energy budget or the gateway's channel capacity.

The problem this thesis addresses is therefore the following: how can identity, integrity and privacy be protected simultaneously in an agricultural data processing




Figure 1.2. *The three protection properties and the shortfall in each. Because all three draw on a single constrained budget, they form one design problem rather than three independent ones.*

pipeline, end to end from constrained sensing node to distributed ledger, within the energy, memory and channel-capacity budgets that field deployment actually permits?

1.3 Rationale of the Study

The rationale rests on an observation about where the budget comes from. If CRT residue decomposition reduces per-packet airtime substantially, it frees channel capacity and node energy that can be reinvested in protection rather than merely banked as efficiency. The same decomposition has a second, less obvious property. An individual residue discloses little about the original value, so residues distributed across channels or custodians provide a natural substrate for confidentiality and for threshold reconstruction. A mechanism introduced for efficiency can thus be made to serve privacy, which is what makes the combination worth studying rather than treating transmission and security as separate layers.

The work is further motivated by an evidential gap. Claims about edge blockchain

feasibility are typically supported by simulation or short laboratory runs. A sustained field deployment on real hardware, through real weather and real connectivity failures, produces a different and more useful class of evidence, including the failure modes that simulation does not generate.

Finally, the applied case is direct. Smallholder and arid-region agriculture is where both the water-efficiency gains and the certification-fraud exposure are largest, and where infrastructure assumptions of existing systems are least likely to hold.

1.4 Objectives of the Study

Objectives are formulated using Bloom’s revised taxonomy, with the cognitive level of each objective stated explicitly so that the verb, the expected evidence and the assessment criterion remain aligned.

1.4.1 Main Objective

To design, implement and evaluate a distributed blockchain-based framework that protects identity, integrity and privacy in agricultural data processing, and to validate it in a sustained field deployment on constrained edge hardware. (*Bloom level: Create*)

1.4.2 Specific Objectives

1. **Analyse** the identity, integrity and privacy requirements of agricultural data pipelines, decomposing the threat surface from sensing node through gateway to ledger into an attacker model with stated capabilities and assumptions. (*Analyse*)
2. **Design** an integrated protection scheme combining CRT multi-channel residue transmission, which lowers per-packet airtime while preserving exact reconstruction,

with an identity layer that binds each reading to an authenticated originating node through HLF MSP credentials and Hierarchical Role-Based Access Control (HRBAC) policy enforced in chaincode. (*Create*)

3. **Implement** the framework on a heterogeneous edge cluster of ESP32 sensing nodes and ARM64 Raspberry Pi (RPi) gateways running HLF under Raft ordering, deployed under real agricultural operating conditions. (*Apply*)
4. **Evaluate** the transmission, ledger and access-control layers against compression and centralised baselines on airtime, energy, throughput, latency and memory footprint, and assess the scalability and fault-tolerance envelope through queuing-theoretic analysis validated against measured load. (*Evaluate*)

1.5 Research Questions

- RQ1.** To what extent can CRT residue decomposition reduce per-packet airtime and increase gateway packet capacity without loss of reconstruction fidelity?
- RQ2.** Can a permissioned distributed ledger sustain the throughput and latency required by real-time irrigation control when hosted entirely on ARM64 edge gateways?
- RQ3.** What identity binding is required for a ledger record to attest to the originating sensing node rather than merely to the submitting gateway, and what does that binding cost?
- RQ4.** What confidentiality is obtained by distributing residues across channels and custodians, and under what adversary model does it hold?
- RQ5.** How does the integrated framework behave as offered load approaches gateway and

ledger saturation, and where are its practical limits?

1.6 Research Hypotheses

H1 (Transmission efficiency). CRT residue transmission reduces per-packet airtime by at least 40% relative to raw transmission while maintaining at least 99% reconstruction accuracy.

H2 (Edge ledger performance). HLF hosted on four RPi 4B gateways sustains a mean transaction latency below 2 s at an average throughput of at least 50 TPS.

H3 (Access-control overhead). Enforcing Role-Based Access Control (RBAC) in chaincode preserves at least 50 TPS and keeps mean smart-contract execution below 500 ms.

H4 (Field robustness). Sustained field operation maintains at least 99% end-to-end reliability, with measured airtime within 5 ms of the analytical model.

H5 (Privacy preservation). An adversary observing fewer than k of the k residue channels cannot reconstruct the original reading, and the residual information leakage remains below a stated bound.

1.7 Scope and Delimitations

In scope. Sub-GHz LoRa sensing networks of the order of fifty nodes; permissioned ledgers of the HLF family under Raft ordering; gateway hardware in the RPi 4B class; agronomic telemetry, specifically soil moisture, temperature, humidity and irrigation actuation; and the identity, integrity and privacy properties of that pipeline.

Out of scope. Public permissionless blockchains and their consensus economics; cellular and satellite backhaul; image-based and remote-sensing agronomy; agronomic model accuracy as such, which is treated as given rather than as an object of study; and formal cryptographic proofs of the underlying primitives, which are cited rather than re-derived.

Delimitations. Field validation is drawn from a single deployment site and a single cropping season, which bounds the claims that can be made about generalisation across soil types and climates. Node population is bounded by what could be instrumented. Behaviour beyond that population is established analytically and by synthetic load rather than by direct measurement.

1.8 Significance of the Study

For **engineering practice**, the work supplies measured rather than simulated evidence about what a permissioned ledger costs on edge hardware, and a transmission scheme that reduces that cost enough to make protection affordable.

For **the research literature**, it connects two bodies of work that have developed separately, namely residue arithmetic for communication efficiency and distributed ledgers for data provenance, and shows that the first can subsidise the second.

For **agricultural stakeholders**, it offers a path to verifiable certification and insurance records that does not require farmers to surrender commercially sensitive detail to every consortium member.

1.9 Definition of Key Terms

Airtime The time a radio occupies a channel to transmit one packet. It governs both node energy consumption and gateway capacity.

Residue The remainder of a reading modulo one of a set of pairwise coprime moduli. A complete residue set uniquely determines the original value below the product of the moduli.

Integrity The property that a record has not been altered since it was committed, and that alteration would be detectable.

Identity The property that a record can be attributed to a specific authenticated principal, here a specific sensing node rather than merely a gateway.

Privacy The property that a record's content is disclosed only to principals authorised to see it, notwithstanding replication of the ledger.

Endorsement In HLF, the process by which designated peers simulate a transaction and sign the result before ordering.

1.10 Organisation of the Thesis

Figure 1.3 shows how the five chapters relate to the four specific objectives of Section 1.4.2.




Figure 1.3. Organisation of the thesis and the chapter in which each specific objective is discharged.

Chapter One states the background, problem, rationale, objectives, questions and hypotheses. **Chapter Two** reviews the literature on precision agriculture and Internet of Things (IoT) sensing, transmission optimisation, distributed ledgers in agriculture, edge-resident consensus and privacy-preserving data processing, and identifies the gap this work addresses. **Chapter Three** presents the materials and methods: system architecture, CRT formulation, identity and access-control design, hardware implementation, deployment protocol and evaluation methodology. **Chapter Four** reports results and applications across the transmission, ledger, access-control and agricultural-outcome dimensions. **Chapter Five** discusses the findings against the hypotheses, states limitations and threats to validity, and sets out conclusions, recommendations and perspectives for further research.

REFERENCES

- [1] M. Rathi and C. Gomathy. “Smart agriculture resource allocation and energy optimization using bidirectional long short-term memory with ant colony optimization (Bi-LSTM–ACO)”. In: *Frontiers in Communications and Networks* 6 (2025), p. 1587402. doi: 10.3389/frcmn.2025.1587402.
- [2] L. S. Pereira, T. Oweis, and A. Zairi. “Irrigation management under water scarcity”. In: *Agricultural Water Management* 57.3 (2002), pp. 175–206. doi: 10.1016/S0378-3774(02)00075-6.
- [3] G. Dono, L. Giraldo, and S. Severini. “The irrigation of agricultural land in Italy: main economic implications and effects in an ongoing climate change”. In: *Agricultural Water Management* 97.7 (2010), pp. 949–957. doi: 10.1016/j.agwat.2010.01.025.
- [4] E. A. Abioye, M. S. Z. Abidin, M. S. A. Mahmud, et al. “A review on monitoring and advanced control strategies for precision irrigation”. In: *Computers and Electronics in Agriculture* 173 (2020), p. 105441. doi: 10.1016/j.compag.2020.105441.
- [5] J. Haxhibeqiri, E. De Poorter, I. Moerman, and J. Hoebeker. “A Survey of LoRaWAN for IoT: From Technology to Application”. In: *Sensors* 18.11 (2018), p. 3995. doi: 10.3390/s18113995.
- [6] F. Adelantado, X. Vilajosana, P. Tuset-Peiro, B. Martinez, J. Melia-Segui, and T. Watteyne. “Understanding the limits of LoRaWAN”. In: *IEEE Communications Magazine* 55.9 (2017), pp. 34–40. doi: 10.1109/MCOM.2017.1600613.
- [7] V. Vukovic, D. Raca, and M. Vukovic. “An Advanced Energy-Efficient Environmental Monitoring in Precision Agriculture Using LoRa-Based Wireless Sensor Networks”. In: *Sensors* 23.14 (2023), p. 6332. doi: 10.3390/s23146332.
- [8] F. Marcelloni and M. Vecchio. “A simple algorithm for data compression in wireless sensor networks”. In: *IEEE Communications Letters* 12.6 (2008), pp. 411–413. doi: 10.1109/LCOMM.2008.080300.
- [9] F. Marcelloni and M. Vecchio. “An Efficient Lossless Compression Algorithm for Periodic Data in Wireless Sensor Networks”. In: *The Computer Journal* 53.6 (2010), pp. 822–833. doi: 10.1093/comjnl/bxp118.
- [10] H. L. Garner. “The residue number system”. In: *IRE Transactions on Electronic Computers* EC-8.2 (1959), pp. 140–147. doi: 10.1109/TEC.1959.5219515.
- [11] C. Ding, D. Pei, and A. Salomaa. *Chinese Remainder Theorem: Applications in Computing, Coding, Cryptography*. World Scientific Publishing, 1996. doi: 10.1142/2893.
- [12] H. H. Pajooh, M. A. Rashid, F. Alam, and S. Demidenko. “Experimental Performance Analysis of a Scalable Distributed Hyperledger Fabric for a Large-Scale IoT Testbed”. In: *Sensors* 22.13 (2022), p. 4868. doi: 10.3390/s22134868.

- [13] E. Androulaki, A. Barger, V. Bortnikov, et al. “Hyperledger Fabric: A Distributed Operating System for Permissioned Blockchains”. In: *Proceedings of the 13th EuroSys Conference*. 2018, 30:1–30:15. doi: 10.1145/3190508.3190538.
- [14] P. Thakkar, S. Nathan, and B. Vishwanathan. “Performance Benchmarking and Optimizing Hyperledger Fabric Blockchain Platform”. In: *2018 IEEE 26th International Symposium on Modeling, Analysis, and Simulation of Computer and Telecommunication Systems (MASCOTS)*. 2018, pp. 264–276. doi: 10.1109/MASCOTS.2018.00034.