

UNIVERSITY OF BUEA

**FACULTY OF ENGINEERING
AND TECHNOLOGY**

**DEPARTMENT OF
COMPUTER ENGINEERING**

**A DISTRIBUTED BLOCKCHAIN-BASED APPROACH TO
PROTECT IDENTITY, INTEGRITY AND PRIVACY IN
AGRICULTURAL DATA PROCESSING**

By

FORCHA GLEN BELOA
(FE22P049)

M.Eng. in Telecommunications and Networks

A Thesis submitted to the Faculty of Engineering and Technology of the
University of Buea in partial fulfilment of the requirements
for the degree of Doctor of Philosophy (PhD)
in Computer Engineering.

Supervisors

Michael Ekonde SONE, PhD
Associate Professor

Elie FUTE TAGNE, PhD
Associate Professor

August 2026

UNIVERSITY OF BUEA

**FACULTY OF ENGINEERING
AND TECHNOLOGY**

**DEPARTMENT OF
COMPUTER ENGINEERING**

**A DISTRIBUTED BLOCKCHAIN-BASED APPROACH TO
PROTECT IDENTITY, INTEGRITY AND PRIVACY IN
AGRICULTURAL DATA PROCESSING**

By

FORCHA GLEN BELOA
(FE22P049)

M.Eng. in Telecommunications and Networks

A Thesis submitted to the Faculty of Engineering and Technology of the
University of Buea in partial fulfilment of the requirements
for the degree of Doctor of Philosophy (PhD)
in Computer Engineering.

Supervisors

Michael Ekonde SONE, PhD
Associate Professor

Elie FUTE TAGNE, PhD
Associate Professor

August 2026

DEDICATION

To . . .

CERTIFICATION

UNIVERSITY OF BUEA
FACULTY OF ENGINEERING AND TECHNOLOGY

This is to certify that the work described in the thesis entitled **“A Distributed Blockchain-Based Approach to Protect Identity, Integrity and Privacy in Agricultural Data Processing”**

By _____

Was carried out in the Department of Computer Engineering

Under the supervision of

(Name and Signature of Supervisor)

Date _____

ACKNOWLEDGEMENTS

ABSTRACT

TABLE OF CONTENTS

.....	x
1 Introduction	1
1.1 Background to the Study	1
1.2 Statement of the Problem	4
1.3 Rationale of the Study	6
1.4 Objectives of the Study	7
1.4.1 Main Objective	7
1.4.2 Specific Objectives	7
1.5 Research Questions	8
1.6 Research Hypotheses	8
1.7 Scope and Delimitations	9
1.8 Significance of the Study	10
1.9 Definition of Key Terms	10
1.10 Organisation of the Thesis	11
2 Review of the Literature	12
2.1 Historical Overview of the Theory and Research Literature	12
2.2 The Theory and Research Literature Specific to the Topic	13
2.2.1 Precision Agriculture and Internet of Things Sensing	13
2.2.2 Low-Power Wide-Area Network Transmission Optimisation and Energy Modelling	14
2.2.3 Residue Number Systems and the Chinese Remainder Theorem in Communications	17
2.2.4 Blockchain and Distributed Ledger Technology in Agriculture and Supply Chains	19
2.2.5 Edge-Resident Consensus and Constrained-Hardware Ledgers ..	21
2.2.6 Identity, Access Control and Privacy-Preserving Data Processing ..	23
2.3 Research in Cognate Areas Relevant to the Topic	25
2.4 Critique of Current Literature and Summary of Known Results	27
2.5 The Contribution This Study Will Make to the Literature	28
A Worked Reconstruction Example for the Chinese Remainder Theorem . . .	42
B Hierarchical Role-Based Access Control Permission Matrix	45

LIST OF FIGURES

1	Lifecycle of an agricultural measurement. A reading becomes a record that is consumed by parties outside the farm, and the value those parties attach to it is what makes falsification worthwhile.	1
2	The three protection properties and the shortfall in each. Because all three draw on a single constrained budget, they form one design problem rather than three independent ones.	5
3	Three lineages of literature underlying this thesis: agricultural sensing, residue arithmetic, and distributed ledgers. Each developed largely independently for decades before beginning to intersect in the 2020s. . . .	13
4	Taxonomy of transmission-efficiency strategies for constrained Internet of Things (IoT) radios. This thesis occupies the shaded leaf on the right that the reviewed literature leaves open.	16
5	The same reconstruction identity (Eq. 2.5) applied to two different parallel resources. Prior work parallelises processor cycles inside one machine; this thesis parallelises radio sub-channels across a shared medium, which introduces identity and confidentiality questions the prior application does not face.	19
6	Consensus fault models and their message complexity, after [115, 117–119].	22

LIST OF TABLES

2	Representative low-power wide-area transmission-efficiency results	17
3	Representative blockchain-in-agriculture systems by authentication boundary and evaluation setting	21
4	Consensus and edge-ledger literature by hardware class evaluated	23
5	Critique of current literature by thread, and what remains unestablished .	27
6	Permission matrix for agricultural resources across the seven Hierarchical Role-Based Access Control (HRBAC) roles. <i>own</i> = own zone only, <i>zone</i> = assigned zone, <i>all</i> = all zones, N/A = no access. Source: [129].	45

LIST OF ABBREVIATIONS

ABAC	Attribute-Based Access Control
ACK	Acknowledgement
ACID	Atomicity, Consistency, Isolation, Durability
ADC	Analogue-to-Digital Converter
API	Application Programming Interface
ARM64	64-bit Advanced RISC Machine architecture
BATMAN-adv	Better Approach To Mobile Ad-hoc Networking, advanced
BLE	Bluetooth Low Energy
BFT	Byzantine Fault Tolerance
CA	Certificate Authority
CAP	Consistency, Availability, Partition tolerance
CFT	Crash Fault Tolerance
CPU	Central Processing Unit
CRT	Chinese Remainder Theorem
DID	Decentralised Identifier
DLT	Distributed Ledger Technology
DP	Differential Privacy
ECDSA	Elliptic Curve Digital Signature Algorithm
EDP	Energy-Delay Product
ESP32	Espressif 32-bit microcontroller with integrated Wi-Fi and Bluetooth
ETSI	European Telecommunications Standards Institute
GDPR	General Data Protection Regulation
GW	Gateway
HLF	Hyperledger Fabric
HRBAC	Hierarchical Role-Based Access Control
IoT	Internet of Things
IPFS	InterPlanetary File System
JSON	JavaScript Object Notation
LoRa	Long Range (sub-GHz spread-spectrum radio modulation)
LoRaWAN	Long Range Wide Area Network
LPWAN	Low-Power Wide-Area Network
LZW	Lempel-Ziv-Welch compression
MAC	Medium Access Control
MCU	Microcontroller Unit
MSP	Membership Service Provider
PBFT	Practical Byzantine Fault Tolerance
PKI	Public Key Infrastructure
RAM	Random Access Memory
RBAC	Role-Based Access Control
RF	Radio Frequency
RNS	Residue Number System
ROI	Return on Investment

RPi	Raspberry Pi
RTC	Real-Time Clock
SF	Spreading Factor
SHA	Secure Hash Algorithm
SLO	Service Level Objective
SSS	Shamir Secret Sharing
TEE	Trusted Execution Environment
TPS	Transactions Per Second
TPM	Trusted Platform Module
WSN	Wireless Sensor Network
X.509	ITU-T public-key certificate standard
ZKP	Zero-Knowledge Proof


CHAPTER ONE

INTRODUCTION

1.1. Background to the Study

Agriculture is undergoing a transition from experience-driven practice to data-driven decision making. Networks of low-cost sensors now measure soil moisture, canopy temperature, humidity and nutrient status continuously, and the resulting streams feed irrigation controllers, yield models and certification schemes [1, 2]. In water-stressed regions this shift is not a convenience but a necessity. Irrigation accounts for the dominant share of freshwater withdrawal, and climate variability has compressed the margin for error in scheduling decisions [3, 4].

A measurement taken in a field does not stop at the field. It becomes a record, and that record is consumed by parties with interests of their own, as Figure 1 sets out. The value of the record to those parties is exactly what creates the incentive to falsify it, and the further a record travels from the instrument that produced it, the weaker the evidence that it is genuine.



*Each downstream use attaches value to the record,
and therefore creates an incentive to falsify it.*

Figure 1. Lifecycle of an agricultural measurement. A reading becomes a record that is consumed by parties outside the farm, and the value those parties attach to it is what makes falsification worthwhile.

The economics of the sensing layer explain why the problem is hard rather than

merely tedious. A deployment that is worth installing on a smallholding must cost a small fraction of the crop value it protects, which forces the node bill of materials down to a level at which neither a hardware security module nor a generous energy budget is affordable. The same economics rule out mains power and wired backhaul across a field, so the design space collapses onto battery-powered microcontrollers speaking a long-range, low-rate radio. Every protection mechanism proposed in this thesis has to survive inside that envelope, and a mechanism that works only on a mains-powered gateway has solved the easier half of the problem.

The sensing layer of such systems is built from severely constrained devices. An Espressif 32-bit microcontroller with integrated Wi-Fi and Bluetooth (ESP32) node operating on battery and harvested solar power must transmit over kilometre-scale distances using Long Range (sub-GHz spread-spectrum radio modulation) (LoRa), a modulation scheme that trades data rate for range and therefore charges a high airtime cost per packet [5, 6]. Airtime is the binding constraint. It determines energy consumption at the node, and, through duty-cycle regulation and collision probability, it determines how many nodes a single gateway can serve [7]. Conventional responses compress the payload, but general-purpose compressors such as Huffman and Lempel-Ziv-Welch compression (LZW) carry dictionary and buffer overheads that a microcontroller with a few kilobytes of free Random Access Memory (RAM) can ill afford [8, 9].

Two further constraints sharpen this. Sub-GHz IoT bands are governed by a duty-cycle limit, so a node that has just transmitted must stay silent for a period proportional to the airtime it consumed. Airtime is therefore not only an energy cost but a rate limit: it caps how often a node may report, and no increase in battery capacity relaxes it. Second, because LoRa uses unslotted access, packets collide whenever two nodes transmit on the same channel at the same time, and the collision probability grows with the square of

offered traffic in the regime that matters. Halving per-packet occupancy therefore buys more than a proportional reduction in collisions, which is why the mechanism studied here is worth the arithmetic it costs.

The Chinese Remainder Theorem (CRT) offers a different route. Decomposing a reading into residues modulo pairwise coprime moduli yields several short packets in place of one long one [10, 11]. Because a multi-channel concentrator such as the SX1302 demodulates channels concurrently, the short residue packets can be received simultaneously, and the original value is recovered in $O(k)$ operations for k moduli. The practical effect is a reduction in per-packet occupancy rather than in total transmitted volume, which is precisely the quantity that governs gateway capacity. The detailed encoding and reconstruction path is presented in Chapter Three.

Residue arithmetic is not new to distributed systems. The same decomposition has been applied to the throughput problem in permissionless blockchains, where splitting work across residue classes permits parallel processing of operations that would otherwise serialise. That line of work establishes the arithmetic and the parallelisation argument; what it does not address is the constrained radio link, where the payoff is airtime rather than processor cycles, nor the identity and confidentiality questions that arise when the residues cross a shared medium. This thesis takes the arithmetic as established and asks what it is worth on the edge of the network rather than at its centre.

Data that has been collected cheaply must then be trusted. Agricultural records increasingly carry consequences beyond the farm gate. They substantiate organic and fair-trade certification, they underwrite index-based crop insurance, and they inform subsidy allocation. Each of these uses creates an incentive to falsify. A centralised database concentrates that risk in a single administrator, and an unprotected sensor network invites the injection of fabricated readings by any party with a compatible radio. Distributed

Ledger Technology (DLT) addresses the tamper-evidence requirement by binding records into a hash-linked chain replicated across mutually distrusting parties, and permissioned platforms such as Hyperledger Fabric (HLF) add an identity layer through a Membership Service Provider (MSP) and Public Key Infrastructure (PKI), making it feasible to say not merely that a record is unaltered but who produced it [12–14].

Running that machinery at the edge is the difficulty. Consensus, endorsement and world-state maintenance were designed for datacentre hardware, whereas an agricultural deployment must place gateways in the field, on 64-bit Advanced RISC Machine architecture (ARM64) single-board computers, behind intermittent connectivity. Whether a permissioned ledger can meet irrigation-control timing requirements on such hardware has until recently been asserted rather than measured.

Finally, protection is not exhausted by integrity. A soil-moisture trace is also a disclosure. It reveals plot productivity, irrigation strategy and, in aggregate, the commercial position of the farmer. Where records are replicated across a consortium that includes buyers, insurers and cooperatives, the same replication that secures integrity enlarges the audience for private information. Identity, integrity and privacy therefore have to be treated as one design problem rather than three independent ones, which is the position this thesis takes.

1.2. Statement of the Problem

Three properties that ought to hold simultaneously are, in current systems, traded against one another.

Identity is weak at the sensing layer. Constrained nodes seldom carry credentials strong enough to authenticate individual readings. Where authentication exists it typically terminates at the gateway, so the ledger attests to what the gateway submitted rather than to what the sensor measured. The provenance chain has a gap exactly where falsification

is cheapest, at the final hop between an individual node and the record that claims to represent it.

Integrity guarantees do not survive the move to the edge. Permissioned ledgers provide tamper evidence, but their published performance characterisations assume server-class hosts. On ARM64 gateways with constrained memory and intermittent links, it is not established that endorsement and ordering can complete inside the control loop of an irrigation decision, nor how the ledger behaves during network partition and recovery.

Privacy is sacrificed to obtain integrity. Replicating raw agronomic readings to every consortium member is the simplest way to make them verifiable and the worst way to keep them confidential. Existing deployments largely accept this trade, or push confidential data off chain in a way that reintroduces the trusted intermediary the ledger was meant to remove.

The three shortfalls, shown together in Figure 2, are not independent faults to be repaired one at a time. Each follows from the same scarcity, and a remedy applied to one is paid for out of the budget available to the others.




Figure 2. The three protection properties and the shortfall in each. Because all three draw on a single constrained budget, they form one design problem rather than three independent ones.

Underlying all three is a resource problem. The transmission budget that would pay for per-reading signatures, encryption and privacy-preserving encodings is already consumed by the raw payload itself. Any scheme that adds cryptographic protection without first reducing transmission cost will exhaust the node's energy budget or the gateway's channel capacity.

How, then, can identity, integrity and privacy be protected simultaneously in an agricultural data processing pipeline, end to end from constrained sensing node to distributed ledger, within the energy, memory and channel-capacity budgets that field deployment actually permits?

1.3. Rationale of the Study

The budget for protection has to come from somewhere. If CRT residue decomposition reduces per-packet airtime substantially, it frees channel capacity and node energy that can be reinvested in protection rather than merely banked as efficiency. The same decomposition has a second, less obvious property. An individual residue discloses little about the original value, so residues distributed across channels or custodians provide a natural substrate for confidentiality and for threshold reconstruction. A mechanism introduced for efficiency can thus be made to serve privacy, which is what makes the combination worth studying rather than treating transmission and security as separate layers.

The two mechanisms also compose cleanly. Residue decomposition and ledger verification operate on the same object at different points in its life, the first on the reading in flight and the second on the record at rest, and the quantity each is sensitive to is different: the radio cares about occupancy, the ledger about the size and count of committed transactions. A scheme that reduces occupancy without increasing transaction count therefore improves one without taxing the other, which is not generally true of

protection mechanisms and is worth demonstrating rather than assuming.

There is also an evidential gap. Claims about edge blockchain feasibility are typically supported by simulation or short laboratory runs. A sustained field deployment on real hardware, through real weather and real connectivity failures, produces a different and more useful class of evidence, including the failure modes that simulation does not generate.

The applied case is direct. Smallholder and arid-region agriculture is where both the water-efficiency gains and the certification-fraud exposure are largest, and where infrastructure assumptions of existing systems are least likely to hold.

1.4. Objectives of the Study

1.4.1. Main Objective

To design, implement and evaluate a distributed blockchain-based framework that protects identity, integrity and privacy in agricultural data processing, and to validate it in a sustained field deployment on constrained edge hardware.

1.4.2. Specific Objectives

1. **Analyse** the identity, integrity and privacy requirements of agricultural data pipelines, decomposing the threat surface from sensing node through gateway to ledger into an attacker model with stated capabilities and assumptions.
2. **Design** an integrated protection scheme combining CRT multi-channel residue transmission, which lowers per-packet airtime while preserving exact reconstruction, with an identity layer that binds each reading to an authenticated originating node through HLF MSP credentials and HRBAC policy enforced in chaincode.

3. **Implement** the framework on a heterogeneous edge cluster of ESP32 sensing nodes and ARM64 Raspberry Pi (RPi) gateways running HLF under Raft ordering, deployed under real agricultural operating conditions.
4. **Evaluate** the transmission, ledger and access-control layers against compression and centralised baselines on airtime, energy, throughput, latency and memory footprint, and assess the scalability and fault-tolerance envelope through queuing-theoretic analysis validated against measured load.

1.5. Research Questions

- RQ1.** To what extent can CRT residue decomposition reduce per-packet airtime and increase gateway packet capacity without loss of reconstruction fidelity?
- RQ2.** Can a permissioned distributed ledger sustain the throughput and latency required by real-time irrigation control when hosted entirely on ARM64 edge gateways?
- RQ3.** What identity binding is required for a ledger record to attest to the originating sensing node rather than merely to the submitting gateway, and what does that binding cost?
- RQ4.** What confidentiality is obtained by distributing residues across channels and custodians, and under what adversary model does it hold?
- RQ5.** How does the integrated framework behave as offered load approaches gateway and ledger saturation, and where are its practical limits?

1.6. Research Hypotheses

- H1 (Transmission efficiency).** CRT residue transmission reduces per-packet airtime by at least 40% relative to raw transmission while maintaining at least 99% reconstruction

accuracy.

H2 (Edge ledger performance). HLF hosted on four RPi 4B gateways sustains a mean transaction latency below 2 s at an average throughput of at least 50 TPS.

H3 (Access-control overhead). Enforcing Role-Based Access Control (RBAC) in chaincode preserves at least 50 TPS and keeps mean smart-contract execution below 500 ms.

H4 (Field robustness). Sustained field operation maintains at least 99% end-to-end reliability, with measured airtime within 5 ms of the analytical model.

H5 (Privacy preservation). An adversary observing fewer than k of the k residue channels cannot reconstruct the original reading, and the residual information leakage remains below a stated bound.

1.7. Scope and Delimitations

In scope. Sub-GHz LoRa sensing networks of the order of fifty nodes; permissioned ledgers of the HLF family under Raft ordering; gateway hardware in the RPi 4B class; agronomic telemetry, specifically soil moisture, temperature, humidity and irrigation actuation; and the identity, integrity and privacy properties of that pipeline.

Out of scope. Public permissionless blockchains and their consensus economics; cellular and satellite backhaul; image-based and remote-sensing agronomy; agronomic model accuracy as such, which is treated as given rather than as an object of study; and formal cryptographic proofs of the underlying primitives, which are cited rather than re-derived.

Delimitations. Field validation is drawn from a single deployment site and a single cropping season, which bounds the claims that can be made about generalisation across

soil types and climates. Node population is bounded by what could be instrumented. Behaviour beyond that population is established analytically and by synthetic load rather than by direct measurement.

1.8. Significance of the Study

For **engineering practice**, the work supplies measured rather than simulated evidence about what a permissioned ledger costs on edge hardware, and a transmission scheme that reduces that cost enough to make protection affordable.

For **the research literature**, it connects two bodies of work that have developed separately, namely residue arithmetic for communication efficiency and distributed ledgers for data provenance, and shows that the first can subsidise the second.

Methodologically, the work contributes a characterisation obtained under field conditions rather than in a laboratory. Failures that a testbed does not generate, mesh partitions, thermal derating, intermittent backhaul and the recovery behaviour that follows each, are exactly the ones that determine whether such a system is deployable, and they only appear when the system is left running.

For **agricultural stakeholders**, it offers a path to verifiable certification and insurance records that does not require farmers to surrender commercially sensitive detail to every consortium member.

1.9. Definition of Key Terms

Airtime The time a radio occupies a channel to transmit one packet. It governs both node energy consumption and gateway capacity.

Residue The remainder of a reading modulo one of a set of pairwise coprime moduli. A complete residue set uniquely determines the original value below the product of the moduli.

Integrity The property that a record has not been altered since it was committed, and that alteration would be detectable.

Identity The property that a record can be attributed to a specific authenticated principal, here a specific sensing node rather than merely a gateway.

Privacy The property that a record's content is disclosed only to principals authorised to see it, notwithstanding replication of the ledger.

Endorsement In HLF, the process by which designated peers simulate a transaction and sign the result before ordering.

1.10. Organisation of the Thesis

Chapter One states the background, problem, rationale, objectives, questions and hypotheses. **Chapter Two** reviews the literature on precision agriculture and IoT sensing, transmission optimisation, distributed ledgers in agriculture, edge-resident consensus and privacy-preserving data processing, and identifies the gap this work addresses. **Chapter Three** presents the materials and methods: system architecture, CRT formulation, identity and access-control design, hardware implementation, deployment protocol and evaluation methodology. **Chapter Four** reports results and applications across the transmission, ledger, access-control and agricultural-outcome dimensions. **Chapter Five** discusses the findings against the hypotheses, states limitations and threats to validity, and sets out conclusions, recommendations and perspectives for further research.

CHAPTER TWO

REVIEW OF THE LITERATURE

2.1. Historical Overview of the Theory and Research Literature

Three lines of research, developed largely apart from one another, converge on the problem this thesis addresses.


The first is the instrumentation of agriculture itself. Wireless sensor networks for environmental monitoring emerged in the early 2000s as a low-cost alternative to wired instrumentation, with early deployments demonstrating that battery-powered nodes could report soil and microclimate variables over months without maintenance [15, 16]. Soil-moisture networks matured through the 2000s and 2010s into standard tools for irrigation scheduling [17, 18], and the arrival of Low-Power Wide-Area Network (LPWAN) radios after 2015 extended the useful range of such networks from tens of metres to several kilometres, at the cost of a severely constrained airtime budget [6, 19]. Precision agriculture as a discipline absorbed this instrumentation gradually, and by the early 2020s the technical barrier to dense sensing had largely fallen, while the barrier to trustworthy, verifiable use of the resulting data had not [20, 21].

The second lineage is arithmetic rather than agricultural. The Chinese Remainder Theorem states a result on simultaneous congruences known in essentially its modern form since antiquity, formalised for modern computation by Garner's constructive reconstruction procedure [10] and developed since as the mathematical basis of residue number systems for fast, parallel, fault-tolerant arithmetic [11, 22]. That literature has run for decades inside computer arithmetic and cryptography, largely untouched by networking or agriculture, its natural home being hardware multipliers, error-correcting codes and, more recently, the parallel verification of cryptographic signatures [23–25].

The third lineage is distributed ledgers. Blockchain moved from a cryptocurrency

mechanism to a general tamper-evident data structure once enterprises recognised that a hash-linked, replicated log could substitute for a trusted central database among mutually distrusting parties [26, 27]. Permissioned platforms such as Hyperledger Fabric formalised that idea for consortium settings, adding an identity layer through a MSP [13, 28], and agricultural supply chains were an early, natural application because certification and provenance are exactly the tamper-evidence problem the ledger solves [29–31].

Figure 3 places these three lineages on one timeline. They do not meet until the last five years, and even where they meet, as Section 2.2 shows, they meet pairwise rather than as a single design.



The three lineages converge only in the present decade, and, as Section 2.2 shows, pairwise rather than as one design.

Figure 3. Three lineages of literature underlying this thesis: agricultural sensing, residue arithmetic, and distributed ledgers. Each developed largely independently for decades before beginning to intersect in the 2020s.

2.2. The Theory and Research Literature Specific to the Topic

2.2.1. Precision Agriculture and Internet of Things Sensing

Precision agriculture research has consistently found that sensing density and irrigation efficiency move together. IoT soil and canopy sensing networks demonstrated early that automated, sensor-driven scheduling outperforms calendar-based irrigation on both water use and yield [1, 2, 33, 34], a finding replicated across arid and semi-arid contexts where the

water-saving incentive is largest [3, 4, 35]. Reviews of the field converge on the same three limitations. First, decision-support integration remains fragmented: sensor data reaches a dashboard but rarely a closed-loop controller with an auditable rationale [20, 21, 36, 37]. Second, most systems evaluate the analytics rather than the network that feeds them, so machine-learning yield and stress models are reported against clean, simulated input streams rather than the loss and jitter that a field deployment actually produces [38–41]. Third, sensing accuracy in the field degrades in ways a laboratory calibration does not capture, drift, fouling and intermittent failure among them, and few studies report the resulting error budget [42–45]. Later generations of this literature push toward higher automation, describing four- and five-generation smart-farming stacks that assume the identity and integrity of every reading rather than establishing it [46–50]. That assumption is exactly the gap Section 2.2.4 returns to.

2.2.2. Low-Power Wide-Area Network Transmission Optimisation and Energy Modelling

LoRa is the dominant physical layer in this literature because it reaches farm-scale distances, of the order of kilometres in rural terrain, from a battery node at sub-milliwatt receive power [5, 19, 51]. The cost is airtime. The Semtech chirp spread-spectrum design fixes the time to transmit a payload of PL bytes as a function of spreading factor SF, bandwidth BW, coding rate CR and the low-data-rate optimisation flag DE [52, 53]. The symbol period is

$$T_{\text{sym}} = \frac{2^{\text{SF}}}{\text{BW}}, \quad (2.1)$$

the payload requires

$$n_{\text{payload}} = 8 + \max\left(\left\lceil \frac{8 PL - 4 SF + 28 + 16 \text{CRC} - 20 \text{IH}}{4(SF - 2 \text{DE})} \right\rceil (\text{CR} + 4), 0\right) \quad (2.2)$$

symbols, and total airtime is

$$T_{\text{air}} = \underbrace{(n_{\text{preamble}} + 4.25) T_{\text{sym}}}_{\text{preamble}} + n_{\text{payload}} T_{\text{sym}}. \quad (2.3)$$


Equation 2.3 is the quantity every optimisation in this literature ultimately reduces: for a fixed SF and BW, airtime is affine in payload length, so a shorter payload is the only lever available once the modulation is chosen. Regulatory duty-cycle limits then convert that airtime into a rate ceiling rather than merely an energy cost. Under European Telecommunications Standards Institute (ETSI) EN 300 220, a sub-GHz node may occupy its channel for at most a fixed fraction DC_{max} of any observation window [54, 55], so a single reading consumes

$$\text{DC} = \frac{T_{\text{air}}}{T_{\text{window}}} \leq \text{DC}_{\text{max}}, \quad (2.4)$$

and offered channel load across N nodes reporting once per window is $G = N T_{\text{air}}/T_{\text{window}}$ Erlang, which under the pure-ALOHA assumption common to this literature bounds per-packet collision probability at $P_{\text{coll}} = 1 - e^{-2G} \approx 2G$ for small G [6, 53]. Node energy per transmission follows directly, $E_{\text{tx}} = P_{\text{tx}} T_{\text{air}}$, so any reduction in T_{air} is simultaneously an energy saving and a capacity gain, which is why the transmission-optimisation literature and the energy-modelling literature are, in practice, one body of work [7, 56–60].

Two families of technique attack Equation 2.3 from opposite sides, summarised in Figure 4. **Compression** keeps one packet per reading and shrinks PL : general lossless coders such as Huffman and Lempel-Ziv-Welch adapted to sensor streams [8, 61–64]

report 20 to 40% payload reduction but require dictionary or frequency-table state that a microcontroller with a few kilobytes of free RAM can only marginally afford [9, 65], and adaptive sampling reduces transmitted volume by skipping readings, which risks missing the threshold event the network exists to detect [66]. **Decomposition** takes the opposite path: it accepts more packets in exchange for each one being far shorter, splitting a reading across parallel channels rather than compressing it on one. Residue-based decomposition for wireless sensor data has been explored for storage reduction and for wireless sensor network lifetime [67–71], but, as Section 2.2.3 details, not for exploiting a multi-channel concentrator’s ability to demodulate several sub-channels *simultaneously*, which is the specific mechanism this thesis studies. LoRa-specific network-capacity work independently confirms that channel occupancy, not raw byte count, is the binding constraint on gateway-served node population [72–77], which is the same quantity Equation 2.4 fixes.



Compression shrinks the payload on one channel; decomposition trades packet count for per-packet occupancy across several. The literature has developed the left branch thoroughly and the right branch for storage and lifetime, but not yet for concurrent multi-channel reception.

Figure 4. Taxonomy of transmission-efficiency strategies for constrained IoT radios. This thesis occupies the shaded leaf on the right that the reviewed literature leaves open.

Table 2 summarises representative results across both branches on the metric that matters for gateway capacity, per-packet airtime or its energy equivalent, rather than the compression ratio each source reports on its own terms.

Table 2. Representative low-power wide-area transmission-efficiency results

Study	Technique	Reported limitation
Li and Simonot-Lion [61]; Marcelloni and Vecchio [8]	Huffman/LZW sensor compression	Dictionary or frequency-table state exceeds the free RAM of the smallest microcontroller class; ratio degrades on low-entropy sensor streams
Kjellsson et al. [66]	Adaptive, event-driven sampling	Reduces volume, not per-packet occupancy; risks missing threshold events between samples
Vidhya and Brindha [67, 68]	CRT residue image compression	Targets storage footprint, not transmission concurrency; single-channel model
Buttigieg et al. [69]; Pourbigharaz and Yassine [70]	RNS energy reduction in WSNs	Energy accounting only; no gateway-capacity or multi-channel reception model
Ye et al. [71]	RNS-based reliable WSN transmission	Reliability under loss, not concurrent channel exploitation

2.2.3. Residue Number Systems and the Chinese Remainder Theorem in Communications

The theorem itself is unconditional. For pairwise coprime moduli $m_1, \dots, m_k$ and $M = \prod_{i=1}^k m_i$, every integer $x \in [0, M)$ is uniquely determined by its residue vector $(r_1, \dots, r_k)$, $r_i = x \bmod m_i$ [10, 24]. Garner's algorithm reconstructs x from the residues in closed form,

$$x = \left(\sum_{i=1}^k r_i M_i N_i \right) \bmod M, \quad M_i = \frac{M}{m_i}, \quad N_i = M_i^{-1} \bmod m_i, \quad (2.5)$$


exact because $M_i N_i \equiv 1 \pmod{m_i}$ by construction [10, 23], and computable in $O(k)$ modular multiplications regardless of the magnitude of x , which is the property that later work in residue number systems generalised into a programme for fast, carry-free, parallel arithmetic [11, 22]. A second property of Equation 2.5 bears on reliability rather than speed. Goldreich, Ron and Sudan formalised the error-tolerant case, showing that x remains recoverable from a strict subset of residues under bounded corruption, "Chinese remaindering with errors" [78], a result later systematised alongside subset-product

reconstruction in computational number theory [25]. That subset-recovery property is a statement about *reconstruction*, not about *disclosure*: it is tempting to read a residue set as a threshold secret-sharing scheme, exact recovery from a complete or near-complete subset, negligible information from any one share in isolation, but the two constructions are not equivalent. A Shamir-style share is deliberately randomised so that any share short of the threshold is uniformly distributed independent of the secret [25]; a plain CRT residue is not. $r_i = x \bmod m_i$ is a deterministic function of x , so intercepting it narrows the field of possible values by a factor of roughly m_i rather than leaking nothing, and if the plaintext domain is itself small relative to m_i , a single residue can narrow that field to very little uncertainty indeed. Whether residue disclosure is negligible in practice therefore depends on the ratio between the plaintext’s domain and the moduli chosen, not on the reconstruction identity alone, and the reviewed literature has not established that ratio for any deployed system; Section 2.2.6 returns to this distinction, and Chapter One’s hypothesis H5 treats it as open rather than established.

Three application communities have used Equations 2.5 for three different resources. The first treats it as a *storage* problem: residue-based lossless image compression reports 15 to 30% footprint reduction with exact reconstruction [67, 68]. The second treats it as an *energy* problem inside wireless sensor networks, converting a single wide-range reading into several narrow-range residues to shrink the arithmetic and storage cost per node [69–71]. The third, established in the candidate’s own prior work [32], treats it as a *processor-cycle* problem: decomposing the modular exponentiation inside Elliptic Curve Digital Signature Algorithm (ECDSA) signature verification into k independent sub-computations over smaller coprime moduli, each assignable to a separate core, so that Bitcoin transaction validation throughput increases without a change to block size or consensus rules. That result establishes the arithmetic and the parallelisation

argument at the centre of a blockchain node; it does not address the constrained radio link at the edge of a sensor network, where the scarce resource is airtime rather than Central Processing Unit (CPU) cycles, nor the identity and confidentiality questions that arise once residues cross a shared wireless medium rather than staying inside one machine's memory. Figure 5 sets the two pipelines side by side. This thesis takes the arithmetic of Equation 2.5 as established by all three communities and asks what it is worth when the parallel resource is a multi-channel radio concentrator rather than a multi-core processor, which, to the reviewed literature's knowledge, no prior system has done [79].

Prior work [32]:



This thesis:




Figure 5. The same reconstruction identity (Eq. 2.5) applied to two different parallel resources. Prior work parallelises processor cycles inside one machine; this thesis parallelises radio sub-channels across a shared medium, which introduces identity and confidentiality questions the prior application does not face.

2.2.4. Blockchain and Distributed Ledger Technology in Agriculture and Supply Chains

A distributed ledger's tamper evidence rests on hash chaining: each block B_i commits to its predecessor and its own transaction set through

$$h_i = H(h_{i-1} \parallel \text{tx}_i), \quad (2.6)$$

where H is a cryptographic hash function, so that altering any committed tx_j for $j \leq i$ requires recomputing every $h_j, \dots, h_i$ across every replica that holds the chain, a cost that grows with the number of mutually distrusting replicas rather than shrinking [80]. Merkle’s original construction generalises Equation 2.6 to a binary tree of hashes so that membership of a single transaction can be proven in $O(\log n)$ rather than by revealing the whole block, a property later systems use for compact audit proofs [81]. Applied to agriculture, this tamper-evidence property has been used for supply-chain transparency [29, 82–88], farm-data management [89–93], and smart-contract automation of certification and payment logic [94–98]. Cross-organisational data sharing has more recently been framed through self-sovereign identity, letting a farm operator authorise fine-grained access without handing a consortium a standing credential [99], and water-quality and irrigation-specific ledgers report similar tamper-evidence gains [100–104]. The generic IoT-blockchain integration literature that agricultural systems draw their architecture from independently confirms the same design pattern: a lightweight device signs and forwards, a gateway or edge node validates and submits, and the ledger anchors a hash rather than the raw payload [86, 105–108].

What this body of literature does not establish is where the tamper evidence begins. Almost every system surveyed authenticates at the gateway or the enterprise boundary and treats the sensor as an untrusted data source whose values are accepted once they arrive, not a principal whose identity is bound to the record it produced [109–113]. Table 3 makes this concrete across a representative sample: the column that is consistently weak or absent is node-level, rather than gateway-level, identity binding, which is precisely the gap Chapter One’s problem statement (Section 1.2) identifies.


Table 3. Representative blockchain-in-agriculture systems by authentication boundary and evaluation setting

Study	Authentication boundary	Evaluation setting	Reported gap
Tian [29]; Kamilaris et al. [82]	Enterprise consortium	/ Conceptual, no field data	No per-sensor identity; throughput not characterised on constrained hardware
Zhang et al. [92]	Gateway	Simulation	Consensus cost assumed server-class hardware
Amraouy et al. [99]	Farm operator (self-sovereign identity)	Prototype	Access-sharing framework; no edge-hardware performance data
Reyna et al. [86]; Dorri et al. [105]	Generic IoT gateway	Survey architecture	/ Identifies the sensor-to-gateway trust gap without closing it

2.2.5. Edge-Resident Consensus and Constrained-Hardware Ledgers

Running a permissioned ledger on gateway-class hardware raises a question the enterprise blockchain literature rarely poses: whether consensus can complete inside the timing budget of a real-time control loop, on a device with an order of magnitude less memory and CPU throughput than the server hardware that literature assumes [12, 13, 114]. Hyperledger Fabric’s execute-order-validate architecture separates transaction endorsement from ordering, letting an ordering service reach agreement on transaction sequence rather than transaction content, and the Raft ordering service used in this thesis implements crash-fault-tolerant, not Byzantine-fault-tolerant, agreement [115, 116]. The distinction has a direct complexity consequence. Crash Fault Tolerance (CFT) protocols such as Raft reach agreement with $O(n)$ message exchanges per decision and tolerate $f = \lfloor (n - 1)/2 \rfloor$ stopped nodes, whereas classical Byzantine Fault Tolerance (BFT) protocols such as Practical Byzantine Fault Tolerance (PBFT) require $n \geq 3f + 1$ replicas, so they tolerate the smaller $f = \lfloor (n - 1)/3 \rfloor$, at $O(n^2)$ messages per decision, to withstand arbitrary, not merely crash, node behaviour [115, 117]. A five-node cluster therefore tolerates two stopped nodes under Raft but only one Byzantine node under PBFT, a

smaller fault budget bought at a quadratically larger message cost. Linear-complexity BFT variants such as HotStuff close much of that gap for adversarial settings [118], and the state-machine-replication literature more broadly treats the CFT/BFT choice as a trade between message cost and the threat model tolerated [119]. This thesis, deployed on a consortium of trusted, physically secured gateways rather than an open or adversarial network, adopts CFT for that reason, a choice the consensus literature supports for permissioned agricultural settings [14] but rarely evaluates on gateway-class ARM64 hardware in the field, as Table 4 shows.



Message complexity scales with the fault model tolerated, not with hardware class. CFT is adopted here because the deployment is a trusted, physically secured consortium; the reviewed literature rarely reports which choice was measured on ARM64 gateway hardware rather than assumed server-class hardware.

Figure 6. Consensus fault models and their message complexity, after [115, 117–119].

Whichever protocol is chosen, its saturation behaviour under load is standard queueing theory. Treating packet or transaction arrivals as Poisson and service time as deterministic (an M/D/1 queue), mean waiting time before service is

$$W_q = \frac{\rho}{2(1 - \rho)} T_{\text{service}}, \quad \rho = \lambda T_{\text{service}}, \quad (2.7)$$

where λ is the arrival rate and T_{service} the deterministic service time, so that total latency $L = T_{\text{service}} + W_q$ is flat while $\rho \ll 1$ and diverges as $\rho \rightarrow 1$. The scalability literature for both

LoRa gateways and blockchain orderers uses Equation 2.7, or its close relatives, to project a saturation point from a small deployment rather than requiring one at scale [120–124], but, as with the consensus literature above, validation against real field load rather than synthetic benchmark load remains the exception rather than the rule. The Consistency, Availability, Partition tolerance (CAP) theorem underlies the same trade from a different angle: no distributed data store can simultaneously guarantee consistency and availability under network partition [125, 126], and permissioned ledgers on intermittently connected rural gateways sit squarely inside that trade-off, a point the practitioner literature on eventual, Atomicity, Consistency, Isolation, Durability (ACID)-relaxed consistency in distributed systems makes explicit [127].

Table 4. Consensus and edge-ledger literature by hardware class evaluated

Study	Protocol	Hardware class	Reported gap
Ongaro and Ousterhout [115]	Raft (CFT)	Server-class (original)	Correctness and understandability proof; no constrained-hardware evaluation
Castro and Liskov [117]	PBFT (BFT)	Server-class	$O(n^2)$ messaging cost not evaluated against gateway bandwidth
Androulaki et al. [13]	Fabric pipeline	E-O-V Data-centre reference deployment	Throughput figures not reproduced on ARM64
Rahimi et al. [123]; Cui et al. [122]	Fabric + sensor / off-chain storage	Prototype, short-duration	No sustained field deployment; saturation point not field-validated

2.2.6. Identity, Access Control and Privacy-Preserving Data Processing

Two threads meet here: how a permissioned ledger decides who may act, and how it protects what a record discloses to parties who are entitled to see that it exists but not what it contains. On the first thread, RBAC enforced in chaincode is the standard mechanism, with permissions attached to roles rather than individual identities and a MSP resolving each transaction’s submitter to a role through its ITU-T public-key certificate

standard (X.509) certificate [28, 128]. The candidate’s own prior formalisation of this mechanism [129] extends flat RBAC with a role hierarchy expressed as a partial order $\preceq$ over the role set R , where $r_i \preceq r_j$ iff $\text{permissions}(r_i) \subseteq \text{permissions}(r_j)$, so that a user u ’s effective permission set in a session is the union across every role dominated by an active role,

$$\text{eff}(u) = \bigcup_{r \in \text{active}(u)} \bigcup_{r' \preceq r} pa(r'), \quad (2.8)$$

with $pa(r)$ the permissions directly assigned to role r . An access request (u, op, res, t) is then granted only if an active, non-expired role of u carries the requested permission *and* the requester’s zone matches the resource’s zone or the requester holds an administrator-issued cross-zone token,

$$\exists r \in \text{active}(u, t) : (op, res) \in \text{eff}(r) \wedge (\text{zone}(u) = \text{zone}(res) \vee \text{czt}(u, \text{zone}(res), t)), \quad (2.9)$$

which folds zone and temporal constraints directly into the access decision rather than leaving them to an external policy layer [129]. This formalisation follows the general tradition of administrative role hierarchies in access-control theory [130] and shares its enforcement point, the chaincode’s client-identity check, with the wider Fabric access-control literature [128], but it authenticates at the *gateway*, the device holding the X.509 certificate that signs the transaction, not at the individual sensing node whose reading the transaction carries. Certificate revocation latency under this model has been measured directly, with propagation delay across a small gossip-connected peer set on the order of minutes rather than seconds [130], a figure this thesis’s own deployment reproduces independently rather than merely assumes.

On the second thread, privacy in a replicated ledger is structurally difficult: the same replication that gives every consortium member a tamper-evident copy of a

record also gives every member a copy of whatever that record discloses. Off-chain storage with an on-chain hash anchor is the common mitigation, verifiable without being world-readable [131], and self-sovereign identity schemes extend the same principle to access itself, letting a data owner grant and revoke visibility without a standing credential [99]. Neither line of work, however, asks whether the transmission layer itself can contribute to confidentiality before a record ever reaches the ledger. Section 2.2.3 showed why that question is harder than it first looks: reconstruction tolerates a missing residue under Goldreich, Ron and Sudan’s error-tolerant formulation [78], but a plain CRT residue is not a randomised secret share, so subset-recovery tolerance does not, by itself, establish subset-disclosure resistance [25]. What it establishes instead is a precise condition worth testing rather than assuming: disclosure from an intercepted residue is bounded by how far the plaintext’s domain exceeds the missing moduli’s product, so confidentiality is a property of the deployed moduli-to-domain ratio, not of the reconstruction identity in the abstract. The literature reviewed here has developed residue decomposition and threshold secret sharing in separate communities without stating that condition, which is exactly why Chapter One frames it as hypothesis H5 rather than a settled result: the rationale (Section 1.3) proposes residue decomposition as a candidate confidentiality substrate, and H5 requires either a quantified leakage bound for the deployed moduli and sensor ranges, or an added randomisation step, before the substrate can be called established.

2.3. Research in Cognate Areas Relevant to the Topic

Three adjacent bodies of work inform this thesis without belonging to any of the six threads above.

General edge computing and IoT security, beyond agriculture, establishes both the promise and the residual risk of pushing computation to the network edge. Edge

processing improves response latency and keeps sensitive data local [132–134], and this holds across smart-city and vehicular deployments as much as agricultural ones, a distributed vehicular network architecture built on Merkle-tree-verified blocks being one such example [81]. The security side of the same literature catalogues IoT-specific attack surfaces, firmware injection, authentication bypass and physical tampering among them [135–141], and a smaller literature addresses the specific hardware this thesis deploys, cache and memory-region behaviour on the ESP32 family, relevant to firmware-level integrity rather than network protocol design [142–144].

Distributed and parallel computing beyond blockchain supplies the scaling vocabulary used in Section 2.2.5 and in Chapter Four: Amdahl’s law bounds the speedup obtainable from parallelising a fixed-size workload by the fraction that remains sequential, a ceiling this thesis’s strong-scaling analysis of gateway count against transaction throughput invokes directly [145]. Ongoing revocation-latency work in vehicular and other resource-constrained permissioned-blockchain settings similarly treats edge computing as the deployment target rather than the enterprise data centre [130], reinforcing the pattern Section 2.2.5 identifies: constrained-hardware evaluation is the literature’s exception rather than its default.

Applied hydrology and arid-region agronomy ground the thesis’s water-efficiency claims outside the networking and ledger literature proper. Global freshwater withdrawal projections identify irrigation as the dominant consumptive use and the primary lever for closing the gap between supply and demand under population growth [146, 147], and Mediterranean and North African basin studies, including region-specific work on Tunisia [148–151], quantify the scale of the shortfall this thesis’s field site sits inside. Crop-specific water-stress and quality studies for tomato and pepper cultivation, the two crops monitored in this thesis’s deployment, supply

the agronomic baselines against which Chapter Four’s yield and water-use results are compared [152–161].

2.4. Critique of Current Literature and Summary of Known Results

Read across the six threads of Section 2.2, the literature is strong on isolated mechanisms and weak on their combination and their field validation. Table 5 states this critique per thread in the same form the thesis’s own contribution (Section 2.5) will answer it.

Table 5. Critique of current literature by thread, and what remains unestablished

Thread	What the literature establishes	What remains unestablished
Precision agriculture & IoT sensing (§2.2.1)	Sensor-driven scheduling reduces water use; decision-support benefit is replicated across regions	Data-integrity assumptions behind those results are rarely stated or measured; field sensing error budgets are under-reported
LPWAN transmission & energy (§2.2.2)	Compression reduces payload; residue decomposition reduces storage and per-node energy in isolation	No prior system uses residue decomposition for concurrent multi-channel <i>reception</i> , the mechanism that governs gateway capacity
RNS / CRT in communications (§2.2.3)	Reconstruction is exact and $O(k)$; subset recovery tolerates missing residues; parallel decomposition accelerates cryptographic verification	The processor-cycle application and the radio-channel application have not been connected; confidentiality implications of residue disclosure are unexploited
Blockchain / DLT in agriculture (§2.2.4)	Hash-chained, replicated ledgers give tamper evidence for agricultural and supply-chain records	Authentication almost universally terminates at the gateway or enterprise boundary, not the sensing node
Edge-resident consensus (§2.2.5)	CFT/BFT trade-off is well characterised in message complexity; CAP constrains any partitioned deployment	Constrained ARM64 gateway-class evaluation, especially under sustained field load rather than benchmark load, is the exception
Identity, access control & privacy (§2.2.6)	RBAC/MSP identity binding and role hierarchies are formalised and enforceable in chaincode	Node-level (not gateway-level) identity binding is largely absent; transmission-layer confidentiality is not treated as a design resource

Three cross-cutting patterns follow from the table. First, every thread that reports a resource saving, transmission, storage, or processor cycles, reports it in isolation, with no accounting of what that saving could purchase elsewhere in the same pipeline; Section 1.3 of Chapter One turns this observation into the thesis’s central rationale. Second, authentication and consensus evaluation both stop at the gateway, leaving the

sensing node, the point at which falsification is cheapest, outside the trust boundary the literature actually measures. Third, laboratory and simulated evaluation dominates every thread; sustained field deployment, of the kind Chapter Three’s methodology and Chapter Four’s results describe, is reported by a minority of the studies reviewed here, and by none that combine transmission efficiency, edge consensus and access control in a single system.

2.5. The Contribution This Study Will Make to the Literature

Against the gaps summarised in Table 5, this thesis contributes the following.

1. It connects the processor-cycle application of CRT parallel decomposition [32] to the radio-channel application this thesis develops, showing that Equation 2.5 pays for a different resource, airtime rather than CPU cycles, when its parallel targets are a concentrator’s demodulation paths rather than a processor’s cores.
2. It extends gateway-terminated identity binding to the sensing node, using the hierarchical RBAC formalisation of Equations 2.8 and 2.9 [129] together with per-node credentials, so that a ledger record attests to the node that produced a reading rather than only to the gateway that submitted it.
3. It reports CFT consensus performance measured on ARM64 gateway hardware under sustained field load rather than benchmark or simulated load, addressing the field-validation gap Table 4 identifies across the consensus literature.
4. It states, and Chapter Four tests as hypothesis H5, the condition under which residue decomposition can serve as a confidentiality substrate rather than merely a transmission optimisation: that disclosure from an intercepted residue is bounded by the ratio between the plaintext domain and the missing moduli’s product, a condition

the subset-recovery results of [25, 78] make precise but do not by themselves establish for any deployed system, connecting the transmission and privacy threads that Section 2.2.6 shows have developed separately.

5. It reports the combined system, transmission, ledger and access control, characterised together on one deployment, rather than each mechanism validated in isolation, which Section 2.4 identifies as the pattern the reviewed literature does not yet show.

REFERENCES

- [1] M. Rathi and C. Gomathy. “Smart agriculture resource allocation and energy optimization using bidirectional long short-term memory with ant colony optimization (Bi-LSTM–ACO)”. In: *Frontiers in Communications and Networks* 6 (2025), p. 1587402. doi: 10.3389/frcmn.2025.1587402.
- [2] L. S. Pereira, T. Oweis, and A. Zairi. “Irrigation management under water scarcity”. In: *Agricultural Water Management* 57.3 (2002), pp. 175–206. doi: 10.1016/S0378-3774(02)00075-6.
- [3] G. Dono, L. Giraldo, and S. Severini. “The irrigation of agricultural land in Italy: main economic implications and effects in an ongoing climate change”. In: *Agricultural Water Management* 97.7 (2010), pp. 949–957. doi: 10.1016/j.agwat.2010.01.025.
- [4] E. A. Abioye, M. S. Z. Abidin, M. S. A. Mahmud, et al. “A review on monitoring and advanced control strategies for precision irrigation”. In: *Computers and Electronics in Agriculture* 173 (2020), p. 105441. doi: 10.1016/j.compag.2020.105441.
- [5] J. Haxhibeqiri, E. De Poorter, I. Moerman, and J. Hoebeker. “A Survey of LoRaWAN for IoT: From Technology to Application”. In: *Sensors* 18.11 (2018), p. 3995. doi: 10.3390/s18113995.
- [6] F. Adelantado, X. Vilajosana, P. Tuset-Peiro, B. Martinez, J. Melia-Segui, and T. Watteyne. “Understanding the limits of LoRaWAN”. In: *IEEE Communications Magazine* 55.9 (2017), pp. 34–40. doi: 10.1109/MCOM.2017.1600613.
- [7] V. Vukovic, D. Raca, and M. Vukovic. “An Advanced Energy-Efficient Environmental Monitoring in Precision Agriculture Using LoRa-Based Wireless Sensor Networks”. In: *Sensors* 23.14 (2023), p. 6332. doi: 10.3390/s23146332.
- [8] F. Marcelloni and M. Vecchio. “A simple algorithm for data compression in wireless sensor networks”. In: *IEEE Communications Letters* 12.6 (2008), pp. 411–413. doi: 10.1109/LCOMM.2008.080300.
- [9] F. Marcelloni and M. Vecchio. “An Efficient Lossless Compression Algorithm for Periodic Data in Wireless Sensor Networks”. In: *The Computer Journal* 53.6 (2010), pp. 822–833. doi: 10.1093/comjnl/bxp118.
- [10] H. L. Garner. “The residue number system”. In: *IRE Transactions on Electronic Computers* EC-8.2 (1959), pp. 140–147. doi: 10.1109/TEC.1959.5219515.
- [11] C. Ding, D. Pei, and A. Salomaa. *Chinese Remainder Theorem: Applications in Computing, Coding, Cryptography*. World Scientific Publishing, 1996. doi: 10.1142/2893.
- [12] H. H. Pajooh, M. A. Rashid, F. Alam, and S. Demidenko. “Experimental Performance Analysis of a Scalable Distributed Hyperledger Fabric for a Large-Scale IoT Testbed”. In: *Sensors* 22.13 (2022), p. 4868. doi: 10.3390/s22134868.

- [13] E. Androulaki, A. Barger, V. Bortnikov, et al. “Hyperledger Fabric: A Distributed Operating System for Permissioned Blockchains”. In: *Proceedings of the 13th EuroSys Conference*. 2018, 30:1–30:15. doi: 10.1145/3190508.3190538.
- [14] P. Thakkar, S. Nathan, and B. Vishwanathan. “Performance Benchmarking and Optimizing Hyperledger Fabric Blockchain Platform”. In: *2018 IEEE 26th International Symposium on Modeling, Analysis, and Simulation of Computer and Telecommunication Systems (MASCOTS)*. 2018, pp. 264–276. doi: 10.1109/MASCOTS.2018.00034.
- [15] N. Wang, N. Zhang, and M. Wang. “Wireless sensors in agriculture and food industry: Recent development and future perspective”. In: *Computers and Electronics in Agriculture* 50.1 (2006), pp. 1–14. doi: 10.1016/j.compag.2005.09.003.
- [16] F. J. Pierce and P. Nowak. “Aspects of precision agriculture”. In: *Advances in Agronomy* 67 (1999), pp. 1–85. doi: 10.1016/S0065-2113(08)60513-1.
- [17] H. R. Boga, J. A. Huisman, C. Oberdörster, and H. Vereecken. “Evaluation of a low-cost soil water content sensor for wireless network applications”. In: *Journal of Hydrology* 344.1 (2010), pp. 32–42. doi: 10.1016/j.jhydrol.2007.06.032.
- [18] R. G. Evans, W. M. Iversen, and Y. Kim. “Integration of wireless sensor network with variable rate irrigation technology”. In: *Proceedings of the ASABE Annual International Meeting* 12 (2008), pp. 1–13.
- [19] E. R. Raj and D. V. Prasad. “A comprehensive study on LoRa technology for wireless sensor networks”. In: *International Journal of Advanced Research in Computer Science* 8.5 (2017), pp. 1214–1219.
- [20] A. McBratney, B. Whelan, and T. Ancev. “Future directions of precision agriculture”. In: *Precision Agriculture* 6 (2005), pp. 7–23. doi: 10.1007/s11119-005-0681-8.
- [21] G. Aydin and M. J. Hall. “A comprehensive review of IoT and wireless sensor networks in precision agriculture”. In: *IEEE Access* 9 (2021), pp. 157733–157755. doi: 10.1109/ACCESS.2021.3129552.
- [22] B. Parhami. *Computer Arithmetic: Algorithms and Hardware Designs*. 2nd. Oxford University Press, 2010.
- [23] D. E. Knuth. *The art of computer programming, volume 2: Seminumerical algorithms*. 3rd. Addison-Wesley Professional, 1997.
- [24] G. H. Hardy, E. M. Wright, R. Heath-Brown, and J. Silverman. *An introduction to the theory of numbers*. 6th. Oxford University Press, 2008.
- [25] V. Shoup. *A computational introduction to number theory and algebra*. 2nd. Cambridge University Press, 2009.
- [26] K. Christidis and M. Devetsikiotis. “Blockchains and smart contracts for the internet of things”. In: *IEEE Access* 4 (2016), pp. 2292–2303.
- [27] J. Gubbi, R. Buyya, S. Marusic, and M. Palaniswami. “Internet of Things (IoT): A vision, architectural elements, and future directions”. In: *Future generation computer systems* 29.7 (2013), pp. 1645–1660.

- [28] Hyperledger Fabric Documentation. *MSP Identity Validity Rules*. <https://hyperledger-fabric.readthedocs.io/en/release-2.2/msp-identity-validity-rules.html>. Accessed: 2025-03-05. 2022.
- [29] F. Tian. “An Agri-Food Supply Chain Traceability System for China Based on RFID and Blockchain Technology”. In: *2016 13th International Conference on Service Systems and Service Management* (2016), pp. 1–6. doi: 10.1109/ICSSSM.2016.7538424.
- [30] X. Xu, C. Pautasso, L. Zhu, et al. “The blockchain as a software connector”. In: *Proceedings of the 13th IEEE International Conference on Software Architecture*. 2016, pp. 182–191. doi: 10.1109/ICSA.2016.15.
- [31] K. Salah, N. Nizamuddin, R. Jayaraman, and M. Omar. “Blockchain-based soybean traceability in agricultural supply chain”. In: *IEEE Access* 7 (2019), pp. 73295–73305. doi: 10.1109/ACCESS.2019.2918000.
- [32] F. G. Beloa. “Enhancing Bitcoin Scalability through Parallel Programming and Chinese Remainder Theorem Optimization with Applications in Agriculture and Healthcare”. Manuscript, University of Buea. Establishes CRT-based decomposition of ECDSA modular arithmetic for parallel transaction validation. 2025.
- [33] C. Verdouw, B. Tekinerdogan, A. Beulens, and S. Wolfert. “Digital twins in smart farming”. In: *Agricultural Systems* 189 (2021), p. 103046. doi: 10.1016/j.agsy.2020.103046.
- [34] N. Gondchawar and R. S. Kawitkar. “IoT Based Smart Agriculture”. In: *International Journal of Advanced Research in Computer and Communication Engineering* 5.6 (2016), pp. 838–842.
- [35] A. Iglesias and L. Garrote. “Adaptation strategies for agricultural water management under climate change in Europe”. In: *Agricultural Water Management* 155 (2015), pp. 113–124. doi: 10.1016/j.agwat.2015.03.014.
- [36] N. Zhang, M. Wang, and N. Wang. “Precision agriculture: a worldwide overview”. In: *Computers and Electronics in Agriculture* 36.2 (2002), pp. 113–132. doi: 10.1016/S0168-1699(02)00096-0.
- [37] R. Gebbers and V. I. Adamchuk. “Precision agriculture and food security”. In: *Science* 327.5967 (2010), pp. 828–831. doi: 10.1126/science.1183899.
- [38] A. Chlingaryan, S. Sukkarieh, and B. Whelan. “Machine learning approaches for crop yield prediction and nitrogen status estimation in precision agriculture: A review”. In: *Computers and Electronics in Agriculture* 151 (2018), pp. 61–69. doi: 10.1016/j.compag.2018.05.012.
- [39] A. Goap, D. Sharma, A. K. Shukla, and C. Rama Krishna. “An IoT based smart irrigation management system using Machine learning and open source technologies”. In: *Computers and Electronics in Agriculture* 155 (2018), pp. 41–49. doi: 10.1016/j.compag.2018.09.040.
- [40] S. P. Mohanty, D. P. Hughes, and M. Salathé. “Using deep learning for image-based plant disease detection”. In: *Frontiers in Plant Science* 7 (2016), p. 1419. doi: 10.3389/fpls.2016.01419.

- [41] T. Van Klompenburg, A. Kassahun, and C. Catal. “Crop yield prediction using machine learning: A systematic literature review”. In: *Computers and Electronics in Agriculture* 177 (2020), p. 105709. doi: 10.1016/j.compag.2020.105709.
- [42] C. M. P. Vaz, S. Jones, M. Meding, and M. Tuller. “Evaluation of Standard Calibration Functions for Eight Electromagnetic Soil Moisture Sensors”. In: *Vadose Zone Journal* 12.2 (2013), vzt2012.0160. doi: 10.2136/vzt2012.0160.
- [43] A. Molina-Cuevas, W. Quijano-Velasco, and A. Guerrero-González. “Low-Cost IoT-Based Soil Monitoring System: Design and Accuracy Assessment”. In: *Sensors* 21.23 (2021), p. 7964. doi: 10.3390/s21237964.
- [44] A. Botta, W. de Donato, V. Persico, and A. Pescapé. “Integration of Cloud computing and Internet of Things: A survey”. In: *Future Generation Computer Systems* 56 (2016), pp. 684–700. doi: 10.1016/j.future.2015.09.021.
- [45] L. Hang and D.-H. Kim. “Design and Implementation of an Integrated IoT Blockchain Platform for Sensing Data Integrity”. In: *Sensors* 19.10 (2019), p. 2228. doi: 10.3390/s19102228.
- [46] R. N. V. Jagan Mohan, P. S. P. S. Rayanoothala, and R. P. R. Praneetha Sree. “Smart agriculture 5.0: blockchain and reinforcement learning synergy for multicropping optimization and traceable IoT-enabled supply chains”. In: *Frontiers in Blockchain* 9 (2026), p. 1766232. doi: 10.3389/fbloc.2026.1766232.
- [47] X. Wang, Q. Wu, H. Zeng, et al. “Blockchain-Empowered H-CPS Architecture for Smart Agriculture”. In: *Advanced Science* 12.27 (2025), p. 2503102. doi: 10.1002/advs.202503102.
- [48] A. Tzounis, N. Katsoulas, T. Bartzanas, and C. Kittas. “Internet of Things in agriculture, recent advances and future challenges”. In: *Biosystems Engineering* 164 (2017), pp. 31–48. doi: 10.1016/j.biosystemseng.2017.09.007.
- [49] L. García, L. Parra, J. M. Jimenez, J. Lloret, and P. Lorenz. “IoT-Based Smart Irrigation Systems: An Overview on the Recent Trends on Sensors and IoT Systems for Irrigation in Precision Agriculture”. In: *Sensors* 20.4 (2020), p. 1042. doi: 10.3390/s20041042.
- [50] P. Jin, W. Du, and W. Jin. “Efficient productivity prediction model based on edge data compression in smart farms”. In: *Smart Agricultural Technology* 12 (Dec. 2025), p. 101242. doi: 10.1016/j.atech.2025.101242.
- [51] S. Li, J. Li, J. Liu, H. Wang, and S. Zhao. “Performance analysis of LoRa networks for smart agriculture”. In: *IEEE Internet of Things Journal* 7.10 (2020), pp. 10045–10057. doi: 10.1109/JIOT.2020.2999323.
- [52] Semtech Corporation. *LoRa Modem Designer’s Guide*. AN1200.22 Rev 2. 2015.
- [53] M. C. Bor, U. Roedig, T. Voigt, and J. M. Alonso. “Do LoRa Low-Power Wide-Area Networks Scale?” In: *Proceedings of the 19th ACM International Conference on Modeling, Analysis and Simulation of Wireless and Mobile Systems (MSWiM)*. ACM, 2016, pp. 59–67. doi: 10.1145/2988287.2989163.
- [54] T. T. Network. “EU868 LoRaWAN Frequency Plan”. In: *TTN Documentation* (2024).
- [55] L. Alliance. “LoRaWAN Regional Parameters v1.0”. In: *LoRa Alliance Technical Specification* (2021).

- [56] K. C. Barr and K. Asanovic. “Energy consumption and bandwidth in wireless sensor networks”. In: *IEEE Micro* 24.6 (2004), pp. 58–68.
- [57] J. Polastre, J. Hill, and D. Culler. “Versatile low power media access for wireless sensor networks”. In: *Proceedings of the 2nd international conference on Embedded networked sensor systems* (2004), pp. 95–107.
- [58] R. Gonzalez and M. Horowitz. “Energy and delay: A metric for early design space exploration”. In: *IEEE Micro* 20.5 (2000), pp. 56–64. doi: 10.1109/40.877952.
- [59] S. Sanchez, S. Blanc, P. Yuste, and J. J. Serrano. “Energy-efficient data transmission in wireless sensor networks for precision agriculture”. In: *Sensors* 21.4 (2021), p. 1209. doi: 10.3390/s21041209.
- [60] W. R. Heinzelman, A. Chandrakasan, and H. Balakrishnan. “Energy-efficient communication protocol for wireless microsensor networks”. In: *Proceedings of the 33rd Annual Hawaii International Conference on System Sciences*. 2000, pp. 10–20. doi: 10.1109/HICSS.2000.926982.
- [61] S. Li, L. Da Xu, and S. Zhao. “5G Internet of Things: A survey”. In: *Journal of Industrial Information Integration* 10 (2018), pp. 1–9. doi: 10.1016/j.jii.2018.01.005.
- [62] J. Ziv and A. Lempel. “A universal algorithm for sequential data compression”. In: *IEEE Transactions on Information Theory* 23.3 (1977), pp. 337–343.
- [63] D. A. Huffman. “A method for the construction of minimum-redundancy codes”. In: *Proceedings of the IRE* 40.9 (1952), pp. 1098–1101.
- [64] D. Salomon. *Data compression: the complete reference*. Springer Science & Business Media, 2007.
- [65] S. O. Community. *What if Dictionary size in LZW algorithm is full?* Stack Overflow. Answer describing typical 4096-entry dictionary limit. 2016.
- [66] J. Kjellsson and A. Zakhori. “Adaptive sampling for sensor networks”. In: *ACM Transactions on Sensor Networks* 4.3 (2008), pp. 1–39. doi: 10.1145/1362542.1362545.
- [67] P. Vidhya and R. Brindha. “Evaluation and Performance Analysis of Chinese Remainder Theorem and Its Application to Lossless Image Compression”. In: *International Journal of Computer Applications* 183.15 (2021), pp. 17–22. doi: 10.5120/ijca2021921556.
- [68] P. Vidhya and R. Brindha. “Evaluation and Performance Analysis of Chinese Remainder Theorem and Its Application to Lossless Image Compression”. In: *International Journal of Computer Applications* 183.15 (2021), pp. 17–22. doi: 10.5120/ijca2021921556.
- [69] V. Buttigieg, R. A. Farrugia, and C. Zammit-Mangion. “Improving lifetime in wireless sensor networks using the Residue Number System”. In: *IET Wireless Sensor Systems* 4.2 (2014), pp. 73–79. doi: 10.1049/iet-wss.2013.0025.
- [70] F. Pourbigharaz and H. M. Yassine. “Modulo-Free Architecture for Binary to Residue Number System Conversion for Moduli Set $\{2^n, 2^n - 1, 2^{n-1} - 1\}$ ”. In: *1994 IEEE International Symposium on Circuits and Systems (ISCAS)* 2 (1994), pp. 317–320. doi: 10.1109/ISCAS.1994.408967.

- [71] R. Ye, A. Boukerche, H. Wang, X. Zhou, and B. Yan. “RESIDENT: A Reliable Residue Number System-Based Data Transmission Mechanism for Wireless Sensor Networks”. In: *Wireless Networks* 24.2 (2018), pp. 505–518. doi: 10.1007/s11276-016-1357-1.
- [72] M. Centenaro, L. Vangelista, A. Zanella, and M. Zorzi. “Long-range communications in unlicensed bands: The LoRa technology”. In: *IEEE Wireless Communications* 23.5 (2016), pp. 12–18.
- [73] H. Zhang, L. He, F. Di Gioia, D. Choi, A. Elia, and P. Heinemann. “LoRaWAN-based Internet of Things (IoT) system for precision irrigation in plasticulture fresh-market tomato”. In: *Smart Agricultural Technology* 2 (2022), p. 100053. doi: 10.1016/j.atech.2022.100053.
- [74] N. J. Chapungo and O. Postolache. “Experimental Evaluation of LoRaWAN Connectivity Reliability in Remote Rural Areas of Mozambique”. In: *Sensors* 25.19 (2025), p. 6027. doi: 10.3390/s25196027.
- [75] E. Bicomakuba, E. Habineza, Samsuzzaman, M. N. Reza, and S.-O. Chung. “IoT-enabled LoRaWAN gateway for monitoring and predicting spatial environmental parameters in smart greenhouses: A review”. In: *Precision Agriculture Science and Technology* 28 (2025), pp. 28–46. doi: 10.22765/pastj.20250003.
- [76] D. Turlykozhaeva, S. Temesheva, N. Ussipov, and A. Baigaliyeva. “Experimental Performance Comparison of Proactive Routing Protocols in Wireless Mesh Network Using Raspberry Pi 4”. In: *Telecom* 5.4 (2024), pp. 1008–1020. doi: 10.3390/telecom5040051.
- [77] M. Bouaziz, A. Rachedi, and A. Belghith. “A Study of LoRaWAN Protocol Performance for IoT Applications in Smart Agriculture”. In: *Computer Communications* 164 (2020), pp. 148–157. doi: 10.1016/j.comcom.2020.10.009.
- [78] O. Goldreich, D. Ron, and M. Sudan. “Chinese remaindering with errors”. In: *Proceedings of the 31st Annual ACM Symposium on Theory of Computing*. 1999, pp. 225–234. doi: 10.1145/301250.301308.
- [79] Y. Wang, I. Chen, and D. Wang. “A novel data compression method for wireless sensor networks”. In: *Sensors* 18.4 (2018), p. 1057. doi: 10.3390/s18041057.
- [80] R. C. Merkle. “A Digital Signature Based on a Conventional Encryption Function”. In: *Lecture Notes in Computer Science* 293 (1988), pp. 369–378. doi: 10.1007/3-540-48184-2_32.
- [81] P. K. Sharma, S. Y. Moon, and J. H. Park. “Block-VN: A Distributed Blockchain Based Vehicular Network Architecture in Smart City”. In: *JIPS: Journal of Information Processing Systems* 13.1 (2017), pp. 184–195. doi: 10.3745/JIPS.04.0022.
- [82] A. Kamilaris, A. Fonts, and F. X. Prenafeta-Boldó. “The rise of blockchain technology in agriculture and food supply chains”. In: *Trends in Food Science & Technology* 91 (2019), pp. 640–652. doi: 10.1016/j.tifs.2019.07.034.
- [83] T. Bosona and G. Gebresenbet. “The Role of Blockchain Technology in Promoting Traceability Systems in Agri-Food Production and Supply Chains”. In: *Sensors* 23.11 (2023), p. 5342. doi: 10.3390/s23115342.

- [84] J. F. Galvez, J. Mejuto, and J. Simal-Gandara. “Future challenges on the use of blockchain for food traceability analysis”. In: *TrAC Trends in Analytical Chemistry* 107 (2018), pp. 222–232. doi: 10.1016/j.trac.2018.08.011.
- [85] F. Tian. “A Supply Chain Traceability System for Food Safety Based on HACCP, Blockchain, and RFID”. In: *2017 International Conference on Applied System Innovation* (2017), pp. 1645–1648. doi: 10.1109/ICASI.2017.7988632.
- [86] A. Reyna, C. Martín, J. Chen, E. Soler, and M. Díaz. “On Blockchain and Its Integration with IoT: Challenges and Opportunities”. In: *Future Generation Computer Systems* 88 (2018), pp. 173–190. doi: 10.1016/j.future.2018.05.046.
- [87] F. Tian. “A supply chain traceability system for food safety based on HACCP, blockchain & Internet of Things”. In: *2017 International Conference on Service Systems and Service Management* (2017), pp. 1–6.
- [88] A. Kamilaris, A. Fonts, and F. X. Prenafeta-Boldú. “The Rise of Blockchain Technology in Agriculture and Food Supply Chains”. In: *Trends in Food Science & Technology* 91 (2019), pp. 640–652. doi: 10.1016/j.tifs.2019.07.034.
- [89] J. Lin, Z. Shen, A. Zhang, and Y. Chai. “Blockchain and IoT based food traceability system for smart agriculture”. In: *Proceedings of the 3rd International Conference on Crowd Science and Engineering*. 2018, pp. 1–6. doi: 10.1145/3265689.3265692.
- [90] M. A. Ferrag, L. Shu, X. Yang, A. Derhab, and L. Maglaras. “Security and privacy for green IoT-based agriculture: Review, blockchain solutions, and challenges”. In: *IEEE Access* 8 (2020), pp. 32031–32053. doi: 10.1109/ACCESS.2020.2973178.
- [91] A. B. Haque, B. Bhushan, and G. Dhiman. “Conceptualizing smart city applications: Requirements, architecture, security issues, and emerging trends”. In: *Expert Systems* 41.5 (2024), e13382. doi: 10.1111/exsy.13382.
- [92] P. Zhang, D. C. Schmidt, J. White, and A. Dubey. “Blockchain for IoT-based smart agriculture”. In: *Proceedings of the IEEE 10th Annual Information Technology, Electronics and Mobile Communication Conference (IEMCON)*. 2019, pp. 0925–0931. doi: 10.1109/IEMCON.2019.8936228.
- [93] A. A. Aliyu and J. Liu. “Blockchain-Based Smart Farm Security Framework for the Internet of Things”. In: *Sensors (Basel, Switzerland)* 23 (Sept. 2023). doi: 10.3390/s23187992.
- [94] R. Rama, V. Kanchana, N. Sivasankari, and S. Prabha. “A systematic review of blockchain technology for agriculture supply chain management”. In: *Materials Today: Proceedings* 47 (2021), pp. 3148–3154. doi: 10.1016/j.matpr.2021.06.253.
- [95] A. S. Patil, B. A. Tama, Y. Park, and K.-H. Rhee. “A framework for blockchain-based smart farming system using IoT”. In: *IEEE Access* 10 (2022), pp. 6845–6862. doi: 10.1109/ACCESS.2022.3142563.
- [96] F. Tian. “An agri-food supply chain traceability system for China based on RFID & blockchain technology”. In: *2016 International Conference on Service Systems and Service Management* (2016), pp. 1–6. doi: 10.1109/ICSSSM.2016.7538424.
- [97] P. Jin, W. Du, and W. Jin. “Efficient productivity prediction model based on edge data compression in smart farms”. In: *Smart Agricultural Technology* 12 (2025), p. 101242. doi: 10.1016/j.atech.2025.101242.

- [98] F. Tian. “A Supply Chain Traceability System for Food Safety Based on HACCP, Blockchain & Internet of Things”. In: *2017 International Conference on Service Systems and Service Management* (2017), pp. 1–6. doi: 10.1109/ICSSSM.2017.7996119.
- [99] O. Amraouy, M. Benbrahim, and M. N. Kabbaj. “A blockchain- and self-sovereign identity-based collaborative framework for secure and verifiable cross-organizational data sharing in smart irrigation”. In: *Smart Agricultural Technology* 12 (Dec. 2025), p. 101654. doi: 10.1016/j.atech.2025.101654.
- [100] L. T. Hung, S. Kang, and J. Song. “Distributed Control Architecture for Blockchain-Based IoT Water Quality Monitoring”. In: *IEEE Access* 10 (2022), pp. 101528–101539. doi: 10.1109/ACCESS.2022.3208440.
- [101] O. Amraouy, M. Benbrahim, and M. N. Kabbaj. “A blockchain- and self-sovereign identity-based collaborative framework for secure and verifiable cross-organizational data sharing in smart irrigation”. In: *Smart Agricultural Technology* 12 (2025), p. 101654. doi: 10.1016/j.atech.2025.101654.
- [102] Aqeel-ur-Rehman, A. U. Abbasi, N. Islam, and Z. A. Shaikh. “A review of wireless sensors and networks’ applications in agriculture”. In: *Computer Standards & Interfaces* 36.2 (2014), pp. 263–270. doi: 10.1016/j.csi.2011.03.004.
- [103] A. A. Aliyu and J. Liu. “Blockchain-Based Smart Farm Security Framework for the Internet of Things”. In: *Sensors* 23.18 (2023), p. 7992. doi: 10.3390/s23187992.
- [104] M. Frustaci, P. Pace, G. Aloï, and G. Fortino. “Evaluating Critical Security Issues of the IoT World: Present and Future Challenges”. In: *IEEE Internet of Things Journal* 5.4 (2018), pp. 2483–2495. doi: 10.1109/JIOT.2017.2767291.
- [105] A. Dorri, S. S. Kanhere, and R. Jurdak. “Blockchain in internet of things: challenges and solutions”. In: *IEEE Communications Surveys & Tutorials* 19.3 (2017), pp. 1397–1408. doi: 10.1109/COMST.2017.2694144.
- [106] A. Dorri, S. S. Kanhere, R. Jurdak, and P. Gauravaram. “Blockchain for IoT Security and Privacy: The Case Study of a Smart Home”. In: *2017 IEEE International Conference on Pervasive Computing and Communications Workshops (PerCom Workshops)* (2017), pp. 618–623. doi: 10.1109/PERCOMW.2017.7917634.
- [107] R. Xu, J. Chen, E. Blasch, and G. Chen. “Making IoT fly: Blockchain-empowered cooperative air-ground networks”. In: *IEEE Network* 34.6 (2020), pp. 288–295. doi: 10.1109/MNET.011.2000111.
- [108] M. T. Hammi, B. Hammi, P. Bellot, and A. Serhrouchni. “Bubbles of Trust: A decentralized blockchain-based authentication system for IoT”. In: *Computers & Security* 78 (2018), pp. 126–142. doi: 10.1016/j.cose.2018.06.004.
- [109] J. Lin, Z. Shen, A. Zhang, and Y. Chai. “Blockchain and IoT based Food Traceability for Smart Agriculture”. In: *Proceedings of the 3rd International Conference on Crowd Science and Engineering* (2018), pp. 1–6. doi: 10.1145/3265689.3265692.
- [110] M. A. Ferrag, L. Shu, and M. Debbah. “Blockchain Technologies for the Internet of Things: Research Issues and Challenges”. In: *IEEE Internet of Things Journal* 7.10 (2020), pp. 10288–10313. doi: 10.1109/JIOT.2020.2984478.

- [111] A. Kamilaris, A. Fonts, and F. X. Prenafeta-Boldú. “The Rise of Blockchain Technology in Agriculture and Food Supply Chains”. In: *Trends in Food Science & Technology* 91 (2019), pp. 640–652. doi: 10.1016/j.tifs.2019.07.034.
- [112] A. Kamilaris, A. Fonts, and F. X. Prenafeta-Boldú. “The Rise of Blockchain Technology in Agriculture and Food Supply Chains”. In: *Trends in Food Science & Technology* 91 (2019), pp. 640–652. doi: 10.1016/j.tifs.2019.07.034.
- [113] A. A. Aliyu and J. Liu. “Blockchain-Based Smart Farm Security Framework for the Internet of Things”. In: *Sensors* 23.18 (2023), p. 7992. doi: 10.3390/s23187992.
- [114] IBM. *Hyperledger Fabric Overview*. <https://www.ibm.com/docs/en/blockchain-platform>. 2023.
- [115] D. Ongaro and J. Ousterhout. “In search of an understandable consensus algorithm”. In: *USENIX Annual Technical Conference* (2014), pp. 305–319.
- [116] H. F. Team. *Ordering Service Documentation*. https://hyperledger-fabric.readthedocs.io/en/latest/orderer/ordering_service.html. 2024.
- [117] M. Castro and B. Liskov. “Practical Byzantine fault tolerance”. In: *Proceedings of the Third Symposium on Operating Systems Design and Implementation (OSDI)*. 1999, pp. 173–186.
- [118] M. Yin, D. Malkhi, M. K. Reiter, G. G. Gueta, and I. Abraham. “HotStuff: BFT consensus with linearity and responsiveness”. In: *Proceedings of the 2019 ACM Symposium on Principles of Distributed Computing (PODC)*. 2019, pp. 347–356. doi: 10.1145/3293611.3331591.
- [119] A. Bessani, J. Sousa, and E. E. Alchieri. “State machine replication for the masses with BFT-SMART”. In: *IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)* (2014), pp. 355–362. doi: 10.1109/DSN.2014.43.
- [120] E. Androulaki, A. Barger, V. Bortnikov, et al. “Hyperledger Fabric: A Distributed Operating System for Permissioned Blockchains”. In: *Proceedings of the 13th EuroSys Conference*. 2018, 30:1–30:15. doi: 10.1145/3190508.3190538.
- [121] C. Gorenflo, S. Lee, L. Golab, and S. Keshav. “FastFabric: Scaling Hyperledger Fabric to 20,000 Transactions per Second”. In: *International Journal of Network Management* 30.5 (2020), e2099. doi: 10.1002/nem.2099.
- [122] D. Mao, Z. Hao, F. Wang, and H. Li. “Innovative Blockchain-Based Approach for Sustainable and Credible Environment in Food Trade: A Case Study in Shandong Province, China”. In: *Sustainability* 10.9 (2018), p. 3149. doi: 10.3390/su10093149.
- [123] H. Sukhwani, J. M. Martinez, X. Chang, K. S. Trivedi, and A. Rindos. “Performance Modeling of PBFT Consensus Process for Permissioned Blockchain Network (Hyperledger Fabric)”. In: *2017 IEEE 36th Symposium on Reliable Distributed Systems (SRDS)* (2017), pp. 253–255. doi: 10.1109/SRDS.2017.36.
- [124] M. A. Rahman, F. Hamzah, and G. A. Safdar. “Residue Number System-based Approach to Minimise Energy Consumption in Wireless Sensor Networks”. In: *IEEE Access* 10 (2022), pp. 51432–51444. doi: 10.1109/ACCESS.2022.3146851.
- [125] E. A. Brewer. “Towards robust distributed systems”. In: *Proceedings of the 19th Annual ACM Symposium on Principles of Distributed Computing (PODC)*. Invited talk. 2000, p. 7.

- [126] S. Gilbert and N. Lynch. “Brewer’s conjecture and the feasibility of consistent, available, partition-tolerant web services”. In: *ACM SIGACT News* 33.2 (2002), pp. 51–59. doi: 10.1145/564585.564601.
- [127] D. Pritchett. “BASE: An ACID alternative”. In: *Queue* 6.3 (2008), pp. 48–55. doi: 10.1145/1394127.1394128.
- [128] P. Długosz et al. “Hyperledger Fabric Access Control System for IoT Layer in Blockchain-Based Applications”. In: *Sensors* 22.20 (2022), pp. 5981–5995. doi: 10.3390/s22205981.
- [129] F. G. Beloa, M. E. Sone, E. F. Tagne, and A. P. Mvogo Asse. “A Decentralized Hierarchical Role-Based Access Control Framework for Agricultural Edge Internet of Things: Design, Implementation, and 61-Day Field Evaluation”. Manuscript submitted to *Future Generation Computer Systems* (Elsevier). Formalises the seven-role HRBAC partial order, zone and temporal constraints, and cascading enrolment protocol this thesis builds on. 2025.
- [130] Z. Lu, W. Wang, X. Zhang, and L. Chen. “Reducing Revocation Latency in IoV using Edge Computing and Permissioned Blockchain”. In: *2022 IEEE International Conference on Blockchain (Blockchain)*. IEEE. 2022, pp. 345–352. doi: 10.1109/Blockchain55522.2022.00055.
- [131] G. Zyskind, O. Nathan, and A. Pentland. “Decentralizing Privacy: Using Blockchain to Protect Personal Data”. In: *2015 IEEE Security and Privacy Workshops* (2015), pp. 180–184. doi: 10.1109/SPW.2015.27.
- [132] S. K. Sharma and X. Wang. “Live data analytics with collaborative edge and cloud processing in wireless IoT networks”. In: *IEEE Internet of Things Journal* 7.7 (2020), pp. 6071–6084. doi: 10.1109/JIOT.2020.2977767.
- [133] Y. Liu, C. Yang, L. Jiang, S. Xie, and Y. Zhang. “Intelligent edge computing for IoT-based energy management in smart cities”. In: *IEEE Network* 33.2 (2019), pp. 111–117. doi: 10.1109/MNET.2019.1800254.
- [134] X. Wang, Y. Han, V. C. M. Leung, D. Niyato, X. Yan, and X. Chen. “Convergence of edge computing and deep learning: A comprehensive survey”. In: *IEEE Communications Surveys & Tutorials* 22.2 (2021), pp. 869–904. doi: 10.1109/COMST.2020.2970550.
- [135] M. S. Farooq, S. Riaz, A. Abid, K. Abid, and M. A. Naeem. “A Survey on the Role of IoT in Agriculture for the Implementation of Smart Farming”. In: *IEEE Access* 7 (2019), pp. 156237–156271. doi: 10.1109/ACCESS.2019.2949703.
- [136] K. Salah, N. Nizamuddin, R. Jayaraman, and M. Omar. “Blockchain-Based Soybean Traceability in Agricultural Supply Chain”. In: *IEEE Access* 7 (2019), pp. 73295–73305. doi: 10.1109/ACCESS.2019.2918000.
- [137] A. Dorri, S. S. Kanhere, and R. Jurdak. “Blockchain in internet of things: challenges and solutions”. In: *IEEE Communications Surveys & Tutorials* 19.3 (2017), pp. 1397–1408. doi: 10.1109/COMST.2017.2694144.
- [138] R. Kalaimathi, A. Yogeshwaran, and P. Rajeswari. “Identifying False Injection Packet Transmission using SVM and Secure OLSR in Wireless Sensor Networks”. In: *4th International Conference on Distributed Computing and Electrical Circuits and Electronics (ICDCECE)*. IEEE. 2025, pp. 1–7. doi: 10.1109/ICDCECE65353.2025.11035330.

- [139] S. S. Kamble, A. Gunasekaran, and S. A. Gawankar. “Sustainable Industry 4.0 framework: A systematic literature review identifying the current trends and future perspectives”. In: *Process Safety and Environmental Protection* 117 (2018), pp. 408–425. doi: 10.1016/j.psep.2018.05.009.
- [140] F.-J. Ferrandez-Pastor, J.-M. Garcia-Chamizo, M. Nieto-Hidalgo, J. Mora-Pascual, and J. Mora-Martinez. “Developing Ubiquitous Sensor Network Platform Using Internet of Things: Application in Precision Agriculture”. In: *Sensors* 16.7 (2016), p. 1141. doi: 10.3390/s16071141.
- [141] M. Lezoche, J. E. Hernandez, M. d. M. E. A. Diaz, H. Panetto, and J. Kacprzyk. “Agri-food 4.0: A survey of the supply chains and technologies for the future agriculture”. In: *Computers in Industry* 117 (2020), p. 103187. doi: 10.1016/j.compind.2020.103187.
- [142] S. Kumari, R. Singh, and A. Kumar. “Building Trust in IoT Networks: A ROM-Based Key Storage and Cloud Signature Verification Approach”. In: *2025 IEEE International Conference on Internet of Things (iThings)*. IEEE. 2025, pp. 1–8. doi: 10.1109/IOT.2025.11089136.
- [143] C. Jay. *ESP32S3 Cache block diagram and operations*. CSDN Blog. Accessed: 2025-10-03. 2025.
- [144] esp-rs Project. *Cache Management and Memory Barriers: esp-hal Documentation*. DeepWiki. Accessed: 2025-03-08. 2025.
- [145] G. M. Amdahl. “Validity of the single processor approach to achieving large scale computing capabilities”. In: *Proceedings of the April 18-20, 1967, Spring Joint Computer Conference*. ACM, 1967, pp. 483–485. doi: 10.1145/1465482.1465560.
- [146] Food and Agriculture Organization of the United Nations. *The future of food and agriculture: Trends and challenges*. <http://www.fao.org/3/a-i6583e.pdf>. 2017.
- [147] United Nations World Water Development Report. *Water for a Sustainable World*. <https://www.unesco.org/reports/wwdr>. 2023.
- [148] Tunisian Ministry of Agriculture. *National Water Strategy 2050*. Government of Tunisia. 2022.
- [149] E. Fereres and M. A. Soriano. “Deficit irrigation for reducing agricultural water use”. In: *Journal of Experimental Botany* 58.2 (2007), pp. 147–159. doi: 10.1093/jxb/erl165.
- [150] L. Dhaouadi, H. Besser, and F. Wassar. “Agriculture sustainability in arid lands of southern Tunisia: Ecological impacts of irrigation water quality and human practices”. In: *Irrigation and Drainage* 69 (2020), pp. 974–996. doi: 10.1002/ird.2492.
- [151] A. Goap, D. Sharma, A. K. Shukla, and C. Rama Krishna. “An IoT based smart irrigation management system using Machine learning and open source technologies”. In: *Computers and Electronics in Agriculture* 155 (2018), pp. 41–49. doi: 10.1016/j.compag.2018.09.040.
- [152] M. Boukhriss, A. Marzouk, I. Cherif, B. Gaalich, T. Triki, and N. Zoghalmi. “Effects of Water Deficit on Growth, Physiology and Fruit Quality Traits of Two Contrasting Tomato (*Solanum lycopersicum* L.) Genotypes”. In: *Agriculture* 11.10 (2021), p. 927. doi: 10.3390/agriculture11100927.

- [153] M. V. Ibáñez, M. D. Muy-Rangel, G. A. González-Aguilar, and J. B. Heredia. “Capsicum Production and Quality: A Review”. In: *HortScience* 55.9 (2020), pp. 1353–1368. doi: 10.21273/HORTSCI15060-20.
- [154] Food and Agriculture Organization. *Good Agricultural Practices for Greenhouse Tomato Production*. FAO Plant Production and Protection Paper. 2013.
- [155] S. Safeer, P. Ricciardi, A. Romano, and G. Russo. “Blockchain-backed sustainable management of the Italian tomato processing industry”. In: *Agriculture* 14.7 (2024), p. 1120. doi: 10.3390/agriculture14071120.
- [156] X. Wang, J. Liu, and Q. Chen. “An advanced deep learning method for pepper diseases and pests detection”. In: *Plant Methods* 21 (2025), p. 70. doi: 10.1186/s13007-025-01387-4.
- [157] A. Iglesias, L. Garrote, F. Flores, and M. Moneo. “Challenges to Manage the Risk of Water Scarcity and Climate Change in the Mediterranean”. In: *Water Resources Management* 21.5 (2007), pp. 775–788. doi: 10.1007/s11269-006-9111-6.
- [158] S. M. Sezen, A. Yazar, and S. Eker. “Effect of drip and sprinkler irrigation on yield and quality of field grown bell pepper”. In: *Agricultural Water Management* 81.1–2 (2006), pp. 115–131. doi: 10.1016/j.agwat.2005.04.008.
- [159] E. Simonne, D. W. Studstill, R. C. Hochmuth, T. Olczyk, and S. E. Webb. “Drip-irrigated pepper production in Florida”. In: *HortTechnology* 16.2 (2006), pp. 283–290. doi: 10.21273/HORTTECH.16.2.0283.
- [160] C. Patanè and S. L. Cosentino. “Effects of soil water deficit on yield and quality of processing tomato under a Mediterranean climate”. In: *Agricultural Water Management* 97.1 (2010), pp. 131–138. doi: 10.1016/j.agwat.2009.08.021.
- [161] S. Topcu, C. Kirda, Y. Dasgan, et al. “Yield response and N-fertiliser recovery of tomato grown under deficit irrigation”. In: *European Journal of Agronomy* 26.1 (2007), pp. 64–70. doi: 10.1016/j.eja.2006.08.004.

APPENDIX A

WORKED RECONSTRUCTION EXAMPLE FOR THE CHINESE REMAINDER THEOREM

Section 2.2.3 states Garner’s reconstruction identity (Equation 2.5) in general form. This appendix carries one complete numeric instance through every step, including the modular-inverse computations that Chapter Two states as results, so that the arithmetic underlying Chapter Three’s implementation is checkable by hand.

Take the moduli set $\{97, 101, 103\}$, pairwise coprime, with product

$$M = 97 \times 101 \times 103 = 1,009,091.$$

A sensor reading of 25.30°C , scaled by 100 to keep the value integral, gives $x = 2530$.

Its three residues are

$$r_1 = 2530 \bmod 97 = 8,$$

$$r_2 = 2530 \bmod 101 = 5,$$

$$r_3 = 2530 \bmod 103 = 58.$$

Step 1: partial products. For each modulus m_i , $M_i = M/m_i$:

$$M_1 = 101 \times 103 = 10,403,$$

$$M_2 = 97 \times 103 = 9,991,$$

$$M_3 = 97 \times 101 = 9,797.$$

Step 2: modular inverses. Each N_i satisfies $M_i N_i \equiv 1 \pmod{m_i}$. Reducing M_i modulo m_i first keeps the inverse search small. For $m_1 = 97$:

$$M_1 \bmod 97 = 10,403 \bmod 97 = 24,$$

$$24 \times 93 = 2,232 = 23 \times 97 + 1 \Rightarrow N_1 = 93.$$

For $m_2 = 101$:

$$M_2 \bmod 101 = 9,991 \bmod 101 = 93,$$

$$93 \times 63 = 5,859 = 58 \times 101 + 1 \Rightarrow N_2 = 63.$$

For $m_3 = 103$:

$$M_3 \bmod 103 = 9,797 \bmod 103 = 12,$$

$$12 \times 43 = 516 = 5 \times 103 + 1 \Rightarrow N_3 = 43.$$

Each inverse is found by inspection here because the moduli are small; the gateway firmware computes the same quantity with the extended Euclidean algorithm, $O(\log m_i)$ per modulus, a sub-millisecond cost at this moduli size (Chapter Three reports measured decomposition and reconstruction timings).

Step 3: reconstruction. Substituting into Equation 2.5:

$$\begin{aligned} x &= (r_1 M_1 N_1 + r_2 M_2 N_2 + r_3 M_3 N_3) \bmod M \\ &= (8 \times 10,403 \times 93 + 5 \times 9,991 \times 63 + 58 \times 9,797 \times 43) \bmod 1,009,091 \\ &= 2,530, \end{aligned}$$

exact, with zero reconstruction error, recovering 25.30°C after descaling. Losing any one residue, say r_3 , leaves a unique solution modulo $m_1m_2 = 9,797$, which happens to be exact for *this particular* value because $2,530 < 9,797$.

Why that qualifier matters. It does not generalise to every encoded value the firmware produces. The single-sensor value worked above fits comfortably under any two-modulus product, but the deployed encoding in `esp32/main/main.c` packs a soil-moisture reading and a temperature bucket into one integer, $x = \text{soil} \times 201 + \text{temp_bucket}$, precisely so that the full three-modulus budget of Section 2.2.3 is used; with soil scaled to $[0, 4095]$, x reaches up to 823,295, which requires all three moduli, since no pair's product exceeds 10,403. The gateway's decode routine (`gateway/crt_decode.py`) accepts any two residues and returns the smallest non-negative solution modulo their product, documenting that this "matches gateway reassembly for sensor values constrained below the product of the received moduli" rather than guaranteeing correctness in general; for a combined value at or above that product, two-residue decoding returns a different, aliased value rather than signalling failure. This is why the subset-recovery figure the source article reports is a measured 64.1% against a 66.7% theoretical figure for two-of-three recovery, not the 100% that would hold if every encoded value stayed under every two-modulus product; when Chapter Three's own data-recovery table is written, it should carry the same measured figure rather than a general bound. The subset-recovery property of Section 2.2.3 is real and exact for values small enough relative to the moduli in play; it is not, as deployed here, a blanket guarantee, and Chapter Three's treatment of loss resilience should state the recovery rate as measured rather than imply the general bound applies unconditionally.

APPENDIX B

HIERARCHICAL ROLE-BASED ACCESS CONTROL PERMISSION MATRIX

Section 2.2.6 states the access-grant predicate (Equation 2.9) that the seven-role HRBAC model evaluates but does not reproduce the permission assignment itself. Both are part of the candidate’s own prior formalisation [129]; Table 6 reproduces it here so that the effective- permission union in Equation 2.8 can be checked against a concrete $pa(r)$ for every role and resource this thesis’s chaincode governs.

Table 6. Permission matrix for agricultural resources across the seven HRBAC roles. *own* = own zone only, *zone* = assigned zone, *all* = all zones, N/A = no access. Source: [129].

Resource	Sensor	Gateway	Farmer	Agron.	Cert.	Supply	Admin
Sensor data	own	zone	all	all	all	N/A	all
Irrigation control	N/A	zone	full	N/A	N/A	N/A	full
Audit logs	N/A	N/A	N/A	N/A	read	N/A	full
Role assignments	N/A	N/A	N/A	N/A	N/A	N/A	full
Cross-zone tokens	N/A	N/A	N/A	N/A	N/A	N/A	issue

Note: Agron. = Agronomist, Cert. = Certification Body, Supply = Supply Chain Partner, following the source’s own column abbreviations [129].

Two rows are worth reading against Equation 2.9 directly. **Sensor data** is the only resource every role can reach in some capacity, own-zone write for the sensor itself, zone-scoped read for the gateway, and unrestricted read for the roles that sit above FARMER in the partial order of Section 2.2.6; this is what makes the zone predicate $zone(u) = zone(res)$ load- bearing rather than redundant; without it, every role above Gateway would read every zone’s sensor data by default. **Role assignments** and **cross-zone tokens** are reachable only by ADMIN, which is what closes Equation 2.9’s $czt(u, zone(res), t)$ predicate: a cross-zone token can only be minted by the one role whose own access is already unrestricted, so the token mechanism widens a requester’s reach without ever widening the set of principals who can create that widening.